

THE EQUITY EXPRESSION OF THE AI-CAPEX THESIS

The capex paradox

Big Tech will spend perhaps three-quarters of a trillion dollars on artificial intelligence this year. This week, four of the five largest spenders tell investors what they are getting for it. The widening gap between the companies buying the compute and the ones selling it has become the cleanest test yet of whether this is a durable investment boom — or the late, expensive stage of one.

A journalist-style explainer built on Heyokha's *AI Capex Wiki* thesis, refreshed with current public-market and earnings-season context. Figures are reference points from the cited sources as of 27–28 July 2026, not live executable quotes.

Not a bet on “AI.”

Hyperscaler shares are operating businesses with depreciation schedules, power bills, debt and cash-flow math. They can compound the AI thesis—and they can also reveal just how expensive winning it has become.

Think of AI as a gold rush. There are two ways to play one: dig, or sell shovels. The shovel-sellers—NVIDIA and the semiconductor complex—are paid in cash today, up front, before anyone knows which claims pan out. The diggers—the hyperscalers pouring money into data centres—are still deep in the mine, spending now against a payoff that arrives years from now, if the seam is as rich as they believe. That timing difference is the whole story. It is why chip shares and hyperscaler shares, supposedly two ends of the same trade, have begun to move in opposite directions.

THE STORY IN 60 SECONDS

~\$725B

Combined 2026 AI capital spending planned by Amazon, Microsoft, Alphabet and Meta—up roughly 77% from an estimated \$410B in 2025, according to analyst tallies compiled by Statista.

–\$5.9B

Alphabet's Q2 2026 free cash flow—negative for the first time as a public company—after quarterly capex doubled to \$44.9B and it lifted full-year guidance to as much as \$205B (Alphabet Q2 results, 22 July).


\$5.09T

NVIDIA's market value on 27 July, the world's most valuable company: the market is still happy to pay up for the shovel-seller even as it questions the diggers (Capital.com / companiesmarketcap).

WHY THE SHARES MOVE ON MORE THAN THE MODEL

Why the equities move faster than “AI”

A model demo is only half the equation. A hyperscaler's basic scorecard is **free cash flow**: operating cash generated, minus the capital spent to build capacity. For a decade that number was enormous and reliable—it funded buybacks, dividends and the “quality” premium these stocks carried. The AI build-out attacks it directly. Capacity is committed years in advance, the chips and buildings depreciate on a fixed schedule, and the power bill arrives monthly—while the revenue that is supposed to justify it all shows up gradually, if at all.




That mechanism cuts both ways. When monetisation outruns the build, cash compounds spectacularly. When the build outruns monetisation—as Alphabet's first negative free-cash-flow quarter suggests it now can—the same fixed costs do not politely wait. So the right question is never simply, “Is AI real?” It is: *where is the gap between what the compute costs and what it earns, and which direction is it moving?*

THE ESSENTIAL DISTINCTION

NVIDIA sells the shovels and is paid today. The hyperscalers dig the mine and are paid—if at all—years from now. Both are called “the AI trade.” They are not the same risk.

A concrete cash-flow example

Alphabet's Q2 2026 print, reported 22 July, makes the arithmetic tangible. The business is not weak—revenue rose 24% year on year to \$119.8B and Google Cloud grew 82% to \$24.8B, with a reported \$514B of cloud backlog. But capital equipment purchases doubled to \$44.9B in the quarter, and that was enough to turn free cash flow negative for the first time in the company's public history.




Source: Alphabet Q2 2026 results and earnings call, 22 July 2026 (Investing.com, CNBC). Free cash flow is not the same as net income; a single quarter is an illustration of the model, not a forecast.

Now scale that picture. If cloud revenue keeps compounding at 80%-plus and eventually laps the depreciation wave, the operating leverage is a gift. But if capex growth stays ahead of monetisation—Alphabet *raised* its 2026 guidance to \$195–205B from \$180–190B even as cash turned negative—the lever swings the other way, and it does so across four balance sheets at once this week.

The 2026 plot twist: the buyers and the sellers split

For most of this cycle “the Magnificent Seven” traded as one idea. In 2026 that idea has fractured. Capital has rotated toward the picks-and-shovels—the semiconductor complex—and away from the mega-cap buyers whose cash flows now carry the build-out. State Street notes the group's average pairwise correlation has fallen to about 0.27 on a three-month basis, down from a 0.78 peak in mid-2025. They are no longer one trade.



* Approximate, different measures: SOX price return YTD (multiple market reports) vs. Mag 7 relative under-performance vs. the S&P 500 YTD (State Street). Shown to convey the split in leadership, not a precise same-basis comparison.

This is the gap at the heart of the debate. Bulls argue the sellers deserve their re-rating—record semiconductor-ETF inflows (a reported ~\$46B in 2026) reflect genuine, visible demand for compute. Skeptics counter that the market is doing something more pointed: rewarding whoever is *paid in cash now* and punishing whoever is *spending it*, the moment that spending shows up as negative free cash flow. Even the shovel-sellers wobbled into earnings week—AMD and Teradyne fell about 5% and 4% on 27 July as doubts about AI returns resurfaced—a reminder that a 47% year-to-date move prices in a great deal of good news.

A USEFUL MENTAL MODEL

AI has a floor made of real demand and scarce compute. But the ceiling—and the path between—is set by financing, power, and proof that the spending earns its keep.

WHY A CORRECTION NEED NOT BREAK THE THESIS

Why the long-term AI case can stay intact

A de-rating of the buyers does not automatically invalidate a multi-year thesis, any more than a bear market in miners disproves gold. Three supports are doing real work underneath the noise.

1. Demand is visible, not hypothetical

Alphabet's cloud backlog of \$514B and 82% cloud growth are booked commitments, not projections. The compute shortage is a real constraint, not a marketing line.

2. It behaves like a general-purpose technology

Heyokha's Wiki cites 13D's Automation Index up 644.5% since 2015 vs. 234.8% for the S&P 500, and agentic-AI deal value up 81.4% year on year—evidence of a broad, compounding platform shift.

3. The build creates hard-asset demand

Data-centre electricity is projected to rise toward ~4.5% of global generation by 2030. That makes power, uranium, natural gas, the grid and copper investable regardless of when AI software pays off.

The measured case, then, is not “bubble” or “boom” but *sequence*: the technology can be transformative and the current equity prices can still be too high for the buyers, because the market is being asked to fund a decade of spending against a monetisation timetable no one can yet prove. As Princeton's Arvind Narayanan put it in a line Heyokha's thesis leans on, the transformation “probably will happen but more on the time scale of something like a decade... not one or two years.”

Better than the last boom—and riskier

The recovery story is not simply that AI is exciting. It is that the largest spenders enter this cycle with real revenue, real cloud franchises and fortress balance sheets—unlike the cash-burning dot-coms of 1999. “Net basis” is the key phrase, though: capacity only creates durable equity value after depreciation, power, debt service and the next round of chips are paid for. On that measure the picture is genuinely mixed.

WHAT HAS CHANGED	WHAT IT LOOKS LIKE IN PRACTICE	WHY IT CUTS BOTH WAYS VS. THE DOT-COM BOOM
Real revenue underneath	Google Cloud +82% to \$24.8B; Azure and AWS growth reaccelerating on AI workloads.	Better: spending is attached to booked demand, not a landgrab. But the revenue still trails the capex by a wide margin.
Balance-sheet strength	Hyperscalers entered with net cash and vast operating cash flow.	Better cushion—yet Heyokha's Wiki flags Nomura's estimate of hyperscaler free cash flow “collapsing to near zero” in 2026 as capex absorbs it.
Capital intensity	Capex-to-sales has surged—e.g. Microsoft from 5.7% (2014) to ~21%, Meta to ~23% (Wiki, Greed & Fear data).	These ratios approach levels that have historically preceded cyclical pullbacks in other capital-heavy industries.
How it's financed	Increasingly with debt: Morgan Stanley (via Wiki) projects ~\$1.15T of AI-build debt through 2028.	The dot-com bust was mostly equity; a debt-funded build introduces credit risk and forced-selling dynamics if returns disappoint.

There is a useful cross-check here. Michael Burry's estimate, cited in the Wiki, is that hyperscalers will direct roughly 94% of 2027 operating cash flow to AI capex, against a historical norm nearer 40%. For the best-positioned franchise that is an aggressive bet on a real opportunity; for the sector as a whole it is exactly the kind of pro-cyclical commitment that turns a slowdown into a bust.

The forces underneath the spend

Reasons the long case is credible

- **Booked demand.** Multi-hundred-billion cloud backlogs and 80%-plus cloud growth are commitments, not hopes.
- **Genuine scarcity.** Leading-edge GPUs, memory and power remain supply-constrained—capacity cannot be summoned on a quarterly timetable.
- **Platform breadth.** Agentic workloads and enterprise adoption are widening the base beyond consumer chat.
- **Physical spillovers.** The build sustains power, uranium, gas, grid and copper demand independent of software monetisation.

Reasons the “spend equals value” thesis can fail

- **Monetisation lag.** If the payoff takes a decade, today's leadership stocks are discounting it far too early (Chris Wood's 1999–2002 parallel).
- **Demand illusion.** Burry argues much current usage is “tokenmaxxing”—benchmarking and trace-harvesting that compresses once firms move in-house.
- **Broken unit economics.** De Gregorio's work (via Wiki) shows AI's marginal costs breaking the classic software-margin contract.
- **Circular, levered financing.** Vendor-customer investment loops and \$1T-plus of committed spend echo pre-2008 structured credit if returns slip.

What to watch this week—not just the tape

The sector becomes far less mysterious if you follow a short dashboard rather than the share-price ticker. These are the readings that separate a durable franchise from an expensive one—and most of them land in the next 72 hours.

SIGNAL	PLAIN-ENGLISH READING	WHY IT MATTERS
Microsoft & Meta (29 Jul)	Capex guidance alongside Azure growth and Meta's ad monetisation.	The two clearest tests of “are they spending <i>into</i> demand, or ahead of it?”
Amazon & Apple (30 Jul)	AWS reacceleration; Apple's more cautious, cash-generative stance.	Shows whether the capex arms race is universal or selective.
Free-cash-flow direction	Is the Alphabet pattern—capex swamping cash—spreading?	Harder to flatter than revenue; it funds the buybacks and dividends these multiples assume.
Capex-guidance language	Any hint of discipline, phasing or “returns focus.”	A shift from a spending race to a returns race would re-rate the whole group.
FOMC (29 Jul)	Target range 3.50–3.75%; ~62% priced for a hold, with inflation near 4.2%.	Long-duration AI bets are most sensitive to the discount rate; a hawkish hold pressures the richest multiples.

WHAT “THE AI TRADE” SHOULD MEAN THIS WEEK

Judge each company on the gap between what its compute costs and what it earns—not on how large the number after “capex” is. A bigger build is not a better business.

How to think about the menu

The shovel-sellers

NVIDIA and the semiconductor complex (SOX). Paid in cash today, with real scarcity—but cyclical, and now pricing in a 47% YTD move. The clean way to own “demand for compute.”

The buyers

The hyperscalers (MSFT, GOOGL, AMZN, META). Genuine franchises with the capex overhang on their cash flows. The wager that monetisation laps depreciation before the market loses patience.

The physical build

Power, uranium, natural gas, grid equipment, copper. The “picks and shovels” of the build-out—investable even if AI software monetisation is slow. Heyokha’s house tilt.

The five names that define this week — representative, not recommendations

GOOGL

Alphabet

reported 22 Jul
–3% after hours

The template for the paradox: a 24% revenue beat and 82% cloud growth, yet the stock fell as capex guidance rose to \$205B and free cash flow went negative. Strong business, contested price.

WATCH NEXT

Whether cloud growth and backlog conversion outrun the depreciation wave now building.

MSFT

Microsoft

reports
29 Jul

Azure growth is the market’s preferred proof-of-monetisation. But the Wiki notes capex-to-sales near 21% and past Copilot-bundling friction as signs the return is not yet obvious.

WATCH NEXT

Azure acceleration vs. capex step-up; any language on capex phasing or AI gross margins.

META

Meta Platforms

reports
29 Jul

The market has, so far, rewarded Meta’s spend when ad revenue delivers. Capex-to-sales near 23% is the highest of the group; the wager is that AI improves the core ads engine now, not someday.

WATCH NEXT

Whether ad monetisation visibly funds the build—or guidance lifts capex again without it.

AMZN

Amazon

reports
30 Jul

AWS is the original cloud franchise and the largest single capex line. Reacceleration would validate the “spending into demand” case; a miss on AWS margins would amplify the buyers’ overhang.

WATCH NEXT

AWS growth and margin trajectory against a rising infrastructure bill.

NVDA

NVIDIA

~\$5.09T mkt cap
27 Jul

The shovel-seller and the cycle’s largest beneficiary—now the world’s most valuable company. Its order book is the hyperscaler capex line, which makes it both the purest bull expression and the most exposed if the buyers blink.

WATCH NEXT

Hyperscaler capex guidance this week is, in effect, NVIDIA’s forward revenue.

The bottom line

Hyperscalers are not “AI with a balance sheet.” They are a leveraged wager that compute demand stays high enough, long enough, for genuinely strong companies to convert an historic capital build into durable per-share cash flow. The 2026 split has made the shovel-sellers expensive and the buyers contested—and this week's earnings will show which description is fairer.

For the bull case on the *buyers* to mature, three things need to happen together: monetisation must show up in cloud revenue and margins, capex growth must decelerate or free cash flow must stabilise, and the debt-financed build must stay orderly. If one fails, the operating leverage that attracts investors becomes the hazard they were being paid to bear.

Editorial note. This feature explains an investment theme; it is not personalised investment advice or a recommendation to buy or sell any security. Historical returns, index observations and the bull/bear framing draw on views attributed to 13D Research, Greed & Fear and other sources compiled in Heyokha's internal Wiki, and should not be treated as independently verified forecasts. Prices and returns are volatile and cut-off dates are stated where used. Heyokha Brothers and the author may hold positions across the names and themes discussed.

Sources and update notes

Internal (Heyokha workspace): Investment Wiki — *AI Capex* thesis (theme; status active; last reviewed 25 June 2026): bull case (13D Automation Index +644.5% since 2015; agentic-AI deal value +81.4%; data-centre electricity toward ~4.5% of global generation by 2030), bear case (Chris Wood's 1999–2002 internet-bubble parallel; capex-to-sales surge; Michael Burry's “tokenmaxxing” and ~94% of 2027 operating cash flow to capex; Ignacio de Gregorio on broken software-margin economics and circular financing; Nomura FCF “collapsing to near zero”; Morgan Stanley ~\$1.15T AI-build debt through 2028), and the Narayanan “decade, not one or two years” quote.

1. Alphabet Q2 2026 results and earnings call, 22 July 2026 (Investing.com; CNBC): revenue +24% to \$119.8B, operating income +30% to \$40.8B, Google Cloud +82% to \$24.8B, \$514B backlog, quarterly capex \$44.9B, free cash flow ~\$5.9B, FY2026 capex guidance raised to \$195–205B; shares ~-3% after hours.

2. Statista, “Big Tech's AI Spending to Reach \$725 Billion in 2026”; Fast Company, Big Tech capex ranked (2026): combined hyperscaler AI capex ~\$725B, up ~77% from ~\$410B in 2025.

3. Zacks / Yahoo Finance earnings preview: reporting calendar — Microsoft and Meta on 29 July; Apple and Amazon on 30 July 2026.

4. Capital.com and companiesmarketcap NVIDIA market-cap updates (27 July 2026): market value ~\$5.09T; share price ~\$196, world's most valuable company.

5. State Street, “Mind on the Market” (20 July 2026): Mag 7 average pairwise correlation ~0.27 (3-month) vs. 0.78 mid-2025 peak; group lagging the S&P 500 YTD.

6. Market reports on 2026 semiconductor leadership (BigGo Finance; intellectia.ai): PHLX Semiconductor Index ~+47% YTD; record ~\$46B semiconductor-ETF inflows in 2026.

7. Seeking Alpha / CNBC market wrap (27 July 2026): AMD ~-5%, Teradyne ~-4% on renewed AI-return doubts; index levels for 27 July.

8. Morningstar; CME; CBS News FOMC previews (July 2026): target range 3.50–3.75%, ~62% priced for a hold at the 28–29 July meeting, US inflation ~4.2%.

Prepared 28 July 2026 · Heyokha Brothers Market Field Guide · Commentary, not investment advice.