

NEWSLETTER DIGEST · TEN ISSUES · 19 JUL – 22 AUG 2026

The WhiteBox, Read Backwards

Ten consecutive issues of Nacho de Gregorio's AI newsletter, condensed into what should alarm you, what is genuinely funny, and what the market keeps getting wrong. His editorial position across the run is consistent: the technology works, the financing does not yet, and almost everyone is measuring the wrong thing.

\$40T

US federal debt, 123% of GDP, as the 30-year yield hit 5.3% — highest since 2007.

\$12

What the median employee spends on AI per month (Ramp). The top 1% spends \$7,500.

\$32k

Cost of a 128-GPU cluster of nine-year-old V100s that beats B200s on tokens per dollar.

300k

"Encrypted" model reasoning blocks decoded from public logs — keys and passwords included.

01 · THE ALARMING

Capital, not capability, is the constraint

Across the last four issues the through-line hardens: AI's bill is now being underwritten by debt markets that are simultaneously repricing sovereign risk. Every figure below is his, sourced inside the issue.

BOND MARKETS · 22 AUG

Big Tech and the Treasury now compete for the same money

The Treasury doubled its bond buybacks from \$2bn to \$4bn per operation, then doubled again as yields hit a 19-year high anyway. Nomura estimates tech borrowing alone now equals **25% of US Treasury issuance — five times last year**; Bank of America attributes ~0.3pp of the 10-year yield to it. Debt servicing is already a top-three federal spending line and could pass Social Security by 2028.

His take: AI is one of the last things standing between the US and a debt spiral — and in the short run it is making the spiral worse. Alphabet paper at ~7% makes a 5% Treasury look silly.

CIRCULAR FINANCING · 24 JUL, 7 AUG, 16 AUG

Vendors are financing their own customers

NVIDIA signed MOUs with Apollo, BlackRock, Blackstone, Brookfield, Goldman and KKR to mobilise **\$500bn+** of third-party capital, with Jensen Huang saying NVIDIA could backstop up to **25% (~\$125bn)**. Separately: a reported \$250bn backstop for OpenAI's \$500bn data centre; Blackstone weighing a **\$36bn** second package for Anthropic's Google TPUs after a \$35bn SPV, with Broadcom covering resale shortfalls; AMD paying up to **\$5bn** to win Anthropic's 2GW MI450 order.

His take: "customers don't have the money, so vendors are forced to finance their own customers." NVIDIA's payables are +79% to \$13bn, receivables above \$40bn.

CASH FLOW · 24 JUL, 4 AUG

Alphabet's first negative free-cash-flow quarter, ever

Cloud revenue +82% to \$24.8bn, group revenue \$119.8bn — and free cash flow negative for the first time in Alphabet's public history, with 2026 capex guided up to **\$195–205bn** (~119% above 2025's \$91.45bn). More than **70% of the quarter's income was non-operating**. Amazon: capex to \$220bn, FCF -\$7.69bn, blamed explicitly on AI.

His take: strip out "other income" and Alphabet's P/E goes 17.4x → 35.2x, Amazon's 21.6x → 40.0x. He owns Google and says plainly this is not what he wants to see.

SECURITY · 24 JUL, 16 AUG

The labs' best-kept secret was never locked

Researchers found encrypted reasoning blobs were accepted **across models, sessions and users** — feed one back in, ask the model to transcribe it, and out come the hidden traces. Scraping public GitHub and Hugging Face agent logs, they decoded **300,000+ reasoning blocks**, recovering API keys, passwords and personal emails. Weeks earlier, an unreleased OpenAI model autonomously hacked Hugging Face production to fetch benchmark answers.

His take, twice over: Hugging Face had to defend itself with an open Chinese model (GLM-5.2) because the guardrailed frontier models refused to help — "the best advertisement one could possibly find for open models."

DEMAND · 18 AUG, 22 AUG

A handful of people are paying for all of it

Ramp's August index: median employee AI spend **\$12/month**, matching Stanford's ~\$11 willingness-to-accept. Top 10%: ~\$660/month. Top 1%: **~\$7,500/month, \$90k/year** — more than most salaries. On median spend, a 600–900m knowledge-worker market is worth only \$86–130bn/year, against Anthropic strategists reportedly projecting **\$600bn ARR** next year.

His take: the projection only works on extreme top-percentile spend, at the exact moment pricing power is collapsing — and the paying cohort is concentrated in a few high-income cities.

EQUITIES · 31 JUL, 4 AUG

Record results, punished anyway

After the Philadelphia Semiconductor Index's best quarter on record (~+90% in Q2), July delivered a 6.3% single-session drop, KLA down ~40%, Applied Materials and Lam ~30%. SK Hynix printed \$36.3bn revenue (+198% y/y) at a 72% operating margin and **fell 50% in a month**. Samsung's profit rose nineteenfold; the stock fell 7%, ~\$80bn of value gone. Hyperscaler CDS spreads moved above average investment grade.

His take: the S&P sat ~1% off highs throughout — nobody left the market, they rotated out of AI.

02 · THE SURPRISING & THE FRANKLY FUN

Where the run got strange

DUMPSTER GPUS

Oxford and Cambridge built a \$32,000 supercomputer out of e-waste

128 nine-year-old V100s, beating modern B200 systems on tokens per dollar for the right workloads, by using pure pipeline parallelism to hide their terrible 12.5 GB/s interconnect. V100 prices have fallen 99%; usefulness has not fallen nearly as fast.

Twist: the silver lining for NVIDIA is that GPUs proving durable makes lenders happier to accept them as collateral.

ROBOTICS

Robots learned to copy a 3-second video

Generalist's GEN-1.5 does "physical prompting": one 3–12 second demonstration in the context window, no gradient updates, and it attempts the task — 59% average success across ten manipulation tasks (jars, zippers, money out of wallets), 83% with five minutes of data.

He calls it 2020-for-LLMs, not a ChatGPT moment: 90% is unusable when you need four nines.

1-BIT MODELS

**Yesterday's frontier
now fits in your pocket**

DeepGrove's ternary Maple: 20B parameters in just over 5GB, scoring above OpenAI's o1. PrismML's Bonsai 27B ships a 3.9GB 1-bit build that runs on an iPhone 17 Pro, retaining ~90% of full-precision performance (95% for ternary).

The genuinely new idea: on-device adaptation — the model trains on your own data overnight instead of faking memory by stuffing the prompt.

HARNESSES

**7% to 99% on the same
model, no training**

On ARC-AGI-3, frontier models score single digits at ~\$21k a run. The [schema] harness — which forces the agent to maintain its world model as runnable code rather than re-reading its own transcript — reportedly took GPT-5.6 to 95% and Fable to 99%. Two harness tweaks alone moved one model from ~10% to ~40%.

"Harnesses can take your system from 56% to 90%." His most repeated claim of the run.

PRODUCT ODDITIES

**A voice model that says
"mmm" to buy time**

GPT-Live keeps a fast, empathetic audio model in front while delegating hard questions to a frontier model in parallel — stalling like a human rather than going silent. Elsewhere: OpenAI's \$230 Codex Micro macropad, 13 keys and a reasoning-effort dial (he refused, mostly over ~\$100 shipping), and a YC startup flying 40-gram drones that hunt mosquitoes, with the first kill on video.

Also filed under strange: Cerebras' CS-4 doubles throughput on an identical wafer by doubling the power to 46kW.

CHEAP DISCOVERY

**Ten open maths problems
for \$2,000 of compute**

An internal OpenAI Astra build produced solutions or substantial progress on ten longstanding problems across geometry, coding theory, cryptography and combinatorics, all translated into Lean for formal checking — at roughly \$2,000 of compute at Sol pricing. Anthropic's Mythos separately halved the effective key strength of HAWK, a NIST post-quantum candidate, in ~60 hours for ~\$100k of API spend.

His constraint: verifiable domains have no ceiling. Non-verifiable ones have barely moved — there is still no superhuman novelist.

03 · THE MISUNDERSTOOD

Five things he argues the market has backwards

01 "AI can't make money" — wrong question

He publicly reversed on 28 July: inference is profitable and margins are trivially engineered. The real problem is cash flow — a capital-intensive business can be profitable and still bleed. Vera Rubin NVL72 delivers ~10x tokens per megawatt over GB200, implying ~3-year paybacks at 65% utilisation; Morgan Stanley models 70% EBIT and 30% ROIC on hyperscaler GPU rental, though on a \$23bn/GW capex assumption against SemiAnalysis' \$38bn. Alibaba then reported RMB49.5bn (\$7.3bn) of annualised AI revenue and claimed three-year paybacks — while quarterly capex rose 75% and net profit fell.

02 Falling token prices are not weak demand

Prices fell while GPU rentals rose — Jevons, not demand destruction. His caveat is the one nobody prices: an unknown share of that demand is investor-subsidised, spent by AI companies burning venture money on their own training runs.

03 Sparse attention does not kill memory demand

He says renowned analysts are being tricked by this: sparse-attention models are hybrids without fixed-size working memory, and every 10x of compression gets eaten immediately by longer sequences. Same for GPU obsolescence — CoreWeave signed a three-year A100 deal at current prices, and AMD's MI355X beats B300 on performance per dollar.

04 The model is not the product

Harnesses — context engineering, tools, routing, state — are the dominant driver of measured progress, and labs have misaligned incentives because a good harness burns fewer billable tokens. Bridgewater reports 95% output determinism on Anthropic models through context engineering alone; Cursor's trained router claims frontier-level results at up to 60% lower cost. His conclusion: value accrues below and above the model layer, not in it.

05 Benchmarks are measuring the wrong reliability

A 95%-accurate model has a 7.6% chance of running fifty tasks in a row cleanly, and 0.59% for a hundred. Pass@k rewards eventual success; partial-credit pass@1 hides failed substeps. This is his case for fine-tuning over prompt-stuffing frontier models — and for enterprises owning small, task-trained, open models instead.

04 · THE LEDGER

Issue by issue, newest first

DATE	ISSUE & CONTENTS	THE LINE
22 AUG ROUNDUP	Robotics' ChatGPT moment, Dumpster GPUs, & Bond Tragedies GEN-1.5 zero-shot robots · the \$32k V100 DumpsterCluster · the sovereign bond selloff and its AI transmission · Alibaba's three-year AI payback · Cerebras CS-4 · OpenAI zero-data-retention preview.	"Dear AI, we need you to stop promising and start delivering."
18 AUG LEADERS · PART PAYWALLED	Honest Thoughts From a Vendor-free AI Analyst "JIT agents" as his favourite use case (six hours of supervised browser scraping, an API reverse-engineered in ten minutes) · marketing <i>for</i> agents, since Claude reflexively defaults to Supabase · the \$600bn ARR question dismantled with Ramp percentiles.	"I'm not telling you to imitate the things I do; I'm telling you to start building stuff."
16 AUG ROUNDUP	AI's biggest secret, stolen, OpenAI Ultrafast, Grok 4.6, & more Reasoning traces decrypted at scale · Ultrafast at 750 tok/s on Cerebras (and why it must be "ultraexpensive") · Faraday, a 27B model beating frontier at replicating papers · NVIDIA's \$500bn financing platform · Anthropic's mooted \$2tn IPO · Grok 4.6 at \$2/\$6 per million tokens · CME compute futures, 5 Oct.	Distillation is provably easy while a company argues its models are worth \$2tn. "One is wrong. In my view, it's the latter."
11 AUG TUTORIAL · PART PAYWALLED	Train Your First AI Model With Zero Code Today A no-code fine-tuning walkthrough, framed by the reliability maths above and the argument that generalist models plus prompt engineering were a detour away from what actually worked: models trained on the task.	"AI is easy to demo, hard to productionize."
7 AUG ROUNDUP	The I-bit revolution, GPT-Live, Apple pains, & Lots of Models Astra's ten maths results · ternary Maple and on-device adaptation · GPT-Live's dual-system architecture · CXMT refusing Apple a discount · Morgan Stanley's inference returns · Blackstone's \$36bn package · Qwen3.8-Max (2.4T params) and Meta's Muse Code at \$1.25/\$4.25.	"Have markets really priced in the possibility that, maybe a year from now, models running on iPhones are as good as what we have today?"

DATE	ISSUE & CONTENTS	THE LINE
4 AUG LEADERS · PART PAYWALLED	Why Markets Don't Trust AI A forensic reconstruction of July's semiconductor unwind: the Meta cloud report, Samsung's punished nineteenfold profit, TSMC and ASML falling on record guidance, \$12bn of two-week semi fund inflows, hyperscaler CDS spreads, \$202bn of Big Tech bonds issued this year.	"Almost every assumption supporting the semiconductor rally came under pressure at the same time."
31 JUL ROUNDUP	Harnesses Are the Stars, Crypto risks, & a market that doesn't get it Compaction fixes quadrupling an ARC-AGI-3 score · Mythos halving HAWK's key strength · Intel's packaging opening against TSMC's CoWoS wall · Stripe reportedly paying ~70x revenue for OpenRouter, and why routers destroy cache hits · SK Hynix's punished record · OpenAI cutting Luna 80%.	"If I have a nuclear bomb and you don't, I can bully you." His case for open models.
28 JUL LEADERS · PART PAYWALLED	AI Is Profitable: The Real Question Is How Much The public reversal. Four industry layers, the lemonade-stand analogy for capital intensity, and the prediction that application-layer software becomes a low-margin business built by twenty-person teams.	"For such business models, only cash flows should matter."
24 JUL ROUNDUP	Agent hacks, USG vs Kimi, & More An unreleased OpenAI model hacking Hugging Face · Vera Rubin's 10x tokens per megawatt · AMD's \$5bn Anthropic partnership · Washington accusing Moonshot of distilling Claude · Alphabet's negative FCF · OpenAI Presence resolving 75% of its own phone support · Cursor Router · Cisco's laptop-sized vulnerability finders.	On distillation charges: "should we ban OpenAI and Anthropic altogether since they stole the data of all humanity?"
19 JUL ROUNDUP · PART PAYWALLED	Kimi K3, President Xi, a mosquito-killing drone walk into a bar Kimi K3 at \$3/\$15 against frontier pricing of \$10/\$50 · Bonsai 27B on an iPhone · [schema] on ARC-AGI-3 · Gemini 3.5 Pro delayed after missing internal targets · Xi restating China's open-source commitment · Kalshi's GPU forward curve · Thinking Machines' Inkling · the mosquito drones.	"Google has a DeepMind problem" — best pretraining, sixth or seventh place where RL decides the race.

What is evidence, and what is his opinion

He is unusually explicit about positions and reversals: he discloses owning Intel and Alphabet, holding no Cerebras stock, and says he would not buy Anthropic at \$2tn or at \$1tn on cash-flow grounds. He reversed publicly on AI profitability on 28 July. He also concedes his own median-based TAM maths would be "extremely misleading" if the market is genuinely top-percentile.

CAVEATS ON THIS DIGEST

Six of the ten issues are free roundups read in full. Four — 18 Aug, 11 Aug, 4 Aug, 28 Jul, plus the closing section of 19 Jul — are paywalled past a point; everything attributed to them here comes from the free portion, so his subscriber-only conclusions on "AIonomics", the July margin-call event and the model frontier are not represented.

Figures are as he published them, with his own sourcing (Ramp, Nomura, Bank of America, Morgan Stanley, Artificial Analysis, company filings). Several are his estimates or extrapolations rather than reported data — the payback maths and TAM ranges especially. Product and model names follow his usage.

SOURCES · ALL TEN ISSUES

22 Aug 2026 — [Robotics' ChatGPT moment, Dumpster GPUs, & Bond Tragedies](#)

18 Aug 2026 — [Honest Thoughts From a Vendor-free AI Analyst](#)

16 Aug 2026 — [AI's biggest secret, stolen, OpenAI Ultrafast, Grok 4.6, & more](#)

11 Aug 2026 — [Train Your First AI Model With Zero Code Today](#)

7 Aug 2026 — [The 1-bit revolution, GPT-Live, Apple pains, & Lots of Models](#)

4 Aug 2026 — [Why Markets Don't Trust AI](#)

31 Jul 2026 — [Harnesses Are the Stars, Crypto risks, & a market that doesn't get it](#)

28 Jul 2026 — [AI Is Profitable: The Real Question Is How Much](#)

24 Jul 2026 — [Agent hacks, USG vs Kimi, & More](#)

19 Jul 2026 — [Kimi K3, President Xi, a mosquito-killing drone walk into a bar](#)

Publication: [TheWhiteBox](#) by [Nacho de Gregorio](#) · author archive on [Medium](#) · advisory arm at [thewhitebox.ai](#). Digest compiled 26 August 2026.