

THE EQUITY EXPRESSION OF THE CHINA AI THESIS

Winning Tokens, Losing Money

Chinese laboratories now serve close to half of the tokens routed through the largest open model marketplace, and they do it at a small fraction of American prices. The shares of the companies that build those models have gone the other way. Alibaba trades 42% below its 52-week high; the two listed pure-play labs, Z.ai and MiniMax, sit roughly three-quarters below theirs. The question is no longer whether China can build competitive AI. It is who, if anyone, gets paid for it.

A reframe worth holding onto

The internal view, carried mainly by White Box, is that China has turned AI into a commodity by design: the same playbook as solar panels, batteries and electric vehicles. Cheapness is not a subsidy that will run out but a strategy, built on lower capital cost per gigawatt and on open weights that let every lab build on every other.

If that is right, the thesis is a claim about volume and diffusion. It is not automatically a claim about shareholder returns.

Method. External figures (prices, company results, market-share data, index levels) come from company releases, exchange-listed quotes and named press reports, each dated in the Sources section. Internal views come from the Heyokha research Wiki page *China AI* (last reviewed 14 September 2026) and its source summaries, and are attributed to the named provider. They are views, not forecasts, and have not been independently verified.

Picture a covered market with two noodle stalls side by side. One charges thirty times the other for a bowl that most customers, most days, cannot tell apart. The cheap stall has the queue; it serves nearly half the diners in the hall. But its owner rebuilds the kitchen every three months because a rival down the aisle has found a faster stove, pays the landlord for the pitch and the gas in advance, and has just gone back to his backers for more money. At the end of the month the people with the most in their pockets are the diners, the restaurants that buy his noodles wholesale and resell them at their own prices, and the landlord. The stall with the queue has won the market. It has not yet won a living.

THE STORY IN 60 SECONDS

46.4%

Chinese models' share of tokens routed on OpenRouter, against 35.7% for US models. A year earlier the average was 11%. *CNBC via Yahoo Finance, 8 Jul 2026*

-42%

Alibaba (9988.HK) at HK\$107.70, against a 52-week high of HK\$186.20, even as cloud revenue grew 45%. *StockAnalysis, 28 Sep 2026*

+122.6%

Growth in first-half net profit at chip designer Cambricon, to RMB2.3bn, on revenue up 108%. *SCMP, 7 Aug 2026*

WHAT ACTUALLY HAPPENED

A price war fought from both ends

Two things happened in September, and they should be read together. On 10 September DeepSeek released V4.1-Flash, the smallest model in a new architecture family, and cut API prices again. Its published off-peak rate is \$0.15 per million input tokens and \$0.60 per million output tokens, doubling at peak hours under a scheme introduced in August. A 1-million-token context window comes as standard. Twelve days later, on 22 September, Anthropic released Claude Opus 5.5 at \$4 input and \$20 output, a 20% cut, and OpenAI answered the same day with GPT-6 Sol at \$2/\$10 and GPT-6 Luna at \$0.10/\$0.50, roughly half the rates of their predecessors.

The first release is the one the internal view anticipated. White Box's 14 September issue called DeepSeek "a top 3 AI lab on the planet" and put its pricing at roughly 200 times below Anthropic's September blended average. The second is the one it did not. At the cheap end of the market, an American lab is now listing a model below DeepSeek's peak-hour rate. The price gap that defined the China AI thesis for eighteen months has not closed at the frontier, where Opus 5.5 still charges thirty-three times DeepSeek's off-peak output rate. But it has been attacked from underneath.

Meanwhile the listed Chinese AI complex had a poor month. Z.ai fell 10.5% on 14 September after announcing a \$5bn raise, its second large fundraising in two months, split between about \$2bn of new shares at HK\$714 and about \$3bn of convertible bonds. It fell a further 12.4% in the week to 25 September, as the Hang Seng Index recorded a third straight weekly loss to close at 24,510.09. Alibaba lost 4.2% over that week, MiniMax 4% and Tencent 2.4%. The proximate cause was caution around the Trump–Xi summit in Washington on 23–25 September, which extended the trade truce, created a new bilateral AI dialogue and an incident hotline, and resolved nothing on chips.

The results season behind all this tells the same story from the inside. Alibaba's June quarter showed cloud revenue up 45% to RMB48.4bn and AI-related product revenue growing at triple digits for a twelfth consecutive quarter. The same release showed capital expenditure up 75% to RMB67.7bn, adjusted EBITA down 30%, net income down 75% and a free-cash outflow of RMB44.7bn. Tencent's second quarter was similar in shape: revenue up 11%, capex up 176%, free cash flow negative. The demand is real. So is the bill.


THE ESSENTIAL DISTINCTION

Share of tokens is not share of revenue, and share of revenue is not share of profit. China can lead the first while trailing the second and third, *and the market is currently paying only for the third.*

THE MECHANISM

Following one million tokens through the economics

A token is the unit a language model reads and writes, roughly three-quarters of an English word. A million of them is a long report, or a few minutes of an autonomous coding agent's work. To see why cheap tokens need not mean profitable tokens, follow a single million output tokens from the price list to the point where someone keeps the surplus.



Step 1: DeepSeek API change log and published pricing (10 Sep 2026); SiliconANGLE (22 Sep 2026). Steps 2–3: internal Wiki views from White Box (10 Jun, 28 Aug 2026), not verified. Step 4: Harvey figures are White Box's account of Bloomberg reporting (23 Sep 2026); Cambricon from SCMP (7 Aug 2026).

The logic of step four is the heart of the matter. In ordinary software, an extra user costs almost nothing to serve, which is why the industry grew used to gross margins above 90%. AI breaks that rule, because every request consumes compute in proportion to its length and ambition. When a supplier sells compute-intensive output into a market where a competitor publishes an open-weight model that anyone can host, the price tends towards the cost of the hardware and electricity needed to run it. What remains is captured by whoever owns something scarce: the chips, the customer relationship, or a distribution channel that already makes money on its own.

That is why the Chinese companies most exposed to the theme split into three groups with very different results. Chip designers and foundries, facing demand that exceeds supply, are making record profits. The platforms, Alibaba and Tencent, can fold cheap models into businesses that already earn money from advertising, commerce and games, but they carry the capital bill. The pure-play labs have neither scarce hardware nor a legacy cash machine, and are funding the gap with repeated share issues.

33×

Opus 5.5 output price (\$20 per million) over DeepSeek V4.1-Flash off-peak (\$0.60). About 17 times at DeepSeek's peak rate.

RMB67.7bn

Alibaba capex in the June quarter, up 75%. That is 5.5 times the quarter's RMB12.4bn of AI-related product revenue.

−RMB44.7bn


Alibaba free cash flow in the same quarter, against an outflow of RMB18.8bn a year earlier.

SiliconANGLE (22 Sep 2026); DeepSeek published pricing (Sep 2026); Alibaba Group June Quarter 2026 results (20 Aug 2026). The 5.5-times ratio is this edition's arithmetic.


THE CENTRAL COMPARISON

Where the price sits, and where the market sits

Listed output price, US\$ per million tokens



Share price against 52-week high, per cent below



Bars are illustrative and drawn to a linear scale; listed prices do not adjust for capability, caching discounts or tokens used per task, and are not like-for-like. Prices: SiliconANGLE (22 Sep 2026), DeepSeek published pricing. Drawdowns: this edition's arithmetic from StockAnalysis last price and 52-week high, quote dates 18–28 Sep 2026 (MiniMax HK\$280.40 vs HK\$1,330; Z.ai HK\$780 vs HK\$2,980; Baidu HK\$87.20 vs HK\$161.20; Alibaba HK\$107.70 vs HK\$186.20; Tencent HK\$436.60 vs HK\$683; SMIC HK\$63.35 vs HK\$93.50). The 52-week highs for the two labs come from a single aggregator and could not be cross-checked; a June report by The Bamboo Works gave earlier peaks of HK\$1,993 and HK\$840.

The two charts are the edition's argument in miniature. The top one shows why China is winning volume: a buyer who does not need the frontier can pay a few per cent of the American price. It also shows the new wrinkle, that an American lab has chosen to compete at the bottom. The lower one shows who the market thinks gets paid. Its ordering is almost exactly the reverse of proximity to the model itself. The companies that sell models as their main product have been marked down hardest; the foundry that sells capacity to all of them, least.

A USEFUL MENTAL MODEL

Think of China AI as China solar, a decade on. The country won the industry by driving the price of the product towards the cost of making it. *The world got cheap panels; the lasting profits went to whoever held a bottleneck, and the panel makers spent years in a price war.*

WHAT HAS ACTUALLY SHIFTED

Six changes, each with two edges

WHAT HAS CHANGED	WHAT IT LOOKS LIKE IN THE DATA	WHY IT CUTS BOTH WAYS
Chinese models lead on volume	46.4% of OpenRouter routed tokens vs 35.7% for US models; DeepSeek alone 17.6%, Qwen 13.9% (Jul 2026).	Proves adoption, but OpenRouter skews towards price-sensitive developers. Revenue share is far lower because Chinese prices are far lower.
US labs now cut too	Opus 5.5 down 20%, cached input down 60%; GPT-6 Sol and Luna about 50% below predecessors (22 Sep).	Validates the Chinese price pressure. It also removes the easy "90% of the quality at 10% of the price" pitch at the low end.
Platform cloud is growing fast	Alibaba cloud +45%, AI-related revenue triple-digit growth for a twelfth quarter (Jun quarter).	The growth is bought with capex up 75% and a RMB44.7bn free-cash outflow. Growth and returns are diverging.
Domestic chips are selling	Cambricon H1 revenue +108%, net profit +122.6%; SMIC says demand is far exceeding expectations.	The bottleneck owners are being paid. Their valuations already reflect it: SMIC trades on about 62 times trailing earnings.
Nvidia access reopens, partly	RTX PRO 5500 inference chips expected to ship from late December; H200 imports far below the 75,000 units licensed per firm.	More compute lowers Chinese costs further. It also dilutes the "self-reliance" story that supports domestic-chip valuations.
Labs keep raising money	Z.ai: \$4bn in July, \$5bn in September; Macquarie models losses into 2030.	Capital keeps the models free to use. Each raise dilutes the holders who are meant to be rewarded for it.

Bull and bear, stated fairly

The bull case	The bear case
• Cost advantage is structural. White Box estimates a fully equipped Chinese gigawatt at about \$15bn against about \$50bn in the US, so each Chinese token carries roughly a third of the capital cost. That is not a subsidy that expires. • Adoption is already won. Chinese share of OpenRouter tokens rose from 4.5% in the first half of 2025 to 46.4% by July 2026 (CNBC). White Box expects commodity tokens to be 80% of the total within a few years. • The platforms monetise indirectly. Alibaba's AI-related revenue has grown at triple digits for twelve straight quarters; Tencent's advertising and games absorb AI without charging for tokens at all. • The chip loop is closing. Cambricon's profits doubled; Alibaba unveiled a Zhenwu V900 chip at three times its predecessor's performance and a 20GW capacity target by 2032 (22 Sep). • Valuations have reset. Tencent trades on about 15 times trailing earnings and Alibaba on about 16 times forward (StockAnalysis). Much of the disappointment may already be in the price. • The ecosystem compounds. In White Box's reading, Chinese labs build on each other's open architectures while US labs compete in secret, so improvements spread faster.	• Commoditisation is the thesis, and it destroys margins. The 13D framing is the solar-panel playbook. That playbook produced cheap products and prolonged losses for producers. • The price floor is now contested. GPT-6 Luna at \$0.10/\$0.50 lists below DeepSeek's peak rate. If US labs are willing to fight at the bottom, China's volume lead costs more to defend. • Capex is running ahead of revenue. Alibaba spent RMB67.7bn in one quarter against RMB12.4bn of AI-related revenue, and adjusted EBITA fell 30%. • The labs depend on new equity. Z.ai has raised \$9bn in two months; Jefferies cites customer concentration, uneven compute supply and low switching costs (Invezz, 14 Sep). • Volume data is narrow. OpenRouter is a developer marketplace. It does not capture enterprise contracts in the US and Europe, where procurement and data rules often exclude Chinese models. • Hardware still binds training. White Box itself says Chinese accelerators are proven for inference but "unproven for frontier-level training", and Huawei says production cannot yet meet domestic demand.

On sourcing: the Wiki's bear section is about capability (compute gaps, distillation, benchmark bias), not about equity returns. The bear case above is therefore built mainly from external reporting and from tensions inside the companies' own results, which is what this edition has done.

Where the internal and external views diverge. The internal view treats the price gap as the durable fact: "in the order of 200 times cheaper" than Anthropic's blended price, in White Box's 14 September words, with China holding the intelligence-per-dollar frontier. Eight days later the external record showed OpenAI listing GPT-6 Luna below DeepSeek's peak-hour price, and Anthropic cutting cached input by 60%. Both can be true. Blended averages include expensive frontier use, and list prices say nothing about quality per token. But the internal view expects Chinese labs to set the floor, and the external data shows the floor is now contested. A second divergence runs through markets. The Wiki's providers, 13D and G&F among them, read China's position in AI as strengthening. The market has marked down every listed name closest to the models. This edition leaves both tensions unresolved.

WHAT TO WATCH FROM HERE

The dashboard

SIGNAL	PLAIN-ENGLISH READING	WHY IT MATTERS
Alibaba September-quarter results	Does cloud growth hold above 40% while capex growth slows?	The first test of whether platform AI spending starts to earn a return rather than just growing.
DeepSeek and Qwen pricing responses	Do Chinese labs cut again in answer to GPT-6 Luna, or hold?	A further cut means the floor is being fought for; holding would suggest they have reached cost.
OpenRouter share after 22 September	Does the Chinese share of routed tokens slip as US prices fall?	Tests whether volume leadership was about price alone or about the models themselves.
RTX PRO 5500 shipments	Nvidia inference chips due from late December, about 500,000 a quarter.	Cheaper Chinese serving costs, but a weaker case for domestic chip champions.

Z.ai and MiniMax funding	Any further placements or convertible issues before year-end.	The clearest gauge of whether pure-play labs can fund themselves from operations.
APEC summit in China, November	Both leaders intend to attend; the new AI dialogue has its first chance to produce something.	Export-control relief or tightening moves the cost base of every name in this edition.

POSTURES

Three postures, not one answer

Own the landlord

Express the theme through the owners of scarce capacity, the chip designers and foundries that every lab must rent from, whichever lab wins. **The cost:** the market already agrees. SMIC trades on about 62 times trailing earnings, and Nvidia's return to China would undercut the scarcity that justifies the price.

Own the distribution

Express it through the platforms, which can give cheap models away inside businesses that already make money from advertising, commerce and games. **The cost:** they carry the capital bill. Alibaba's free cash flow was -RMB44.7bn in one quarter, and the return on that spending is unproven.

Own the reset labs

Express it through the listed pure-play labs, now roughly three-quarters below their highs, on the view that volume eventually turns into pricing power. **The cost:** repeated dilution, losses that analysts model into 2030, and a business in which each new release resets the price of the last.

The six companies below are illustrative names chosen to show how different parts of the theme behave. They are NOT recommendations, and nothing in this edition is advice to buy, sell or hold any security.

9988.HK

Alibaba Group

HK\$107.70
28 Sep 2026

The theme's biggest platform. June-quarter cloud revenue rose 45% to RMB48.4bn, and the Qwen family is among the most used open models. It also shows the cost most clearly: capex up 75%, adjusted EBITA down 30%, free cash flow -RMB44.7bn. At Apsara on 22 September it set a 20GW capacity goal for 2032 and unveiled the Zhenwu V900 chip.

WATCH NEXT

Whether September-quarter capex growth slows while cloud growth holds above 40%.

0700.HK

Tencent Holdings

HK\$436.60
25 Sep 2026

The distribution case in its purest form. Second-quarter revenue rose 11% to RMB204.8bn and adjusted net income reached RMB68.4bn, with AI showing up in advertising and games rather than in token sales. Capex rose 176% and free cash flow turned negative. On about 15 times trailing earnings, it offers exposure to AI adoption with little direct exposure to token pricing.

WATCH NEXT

Whether advertising growth keeps pace with the capex ramp in the third quarter.

2513.HK

Z.ai (Zhipu)

HK\$780.00
18 Sep 2026

Developer of the GLM models the Wiki credits as China's first frontier-level release. First-half revenue rose 400% to RMB953.9m. It announced a \$5bn raise on 14 September, at HK\$714 a share and about \$3bn in convertibles, two months after a \$4bn raise. The stock fell 10.5% that day and a further 12.4% in the week to 25 September.

WATCH NEXT

Conversion terms and any further issuance before year-end.

0100.HK

MiniMax Group

HK\$280.40
24 Sep 2026

A multimodal lab (language, video, speech) with consumer products such as Hailuo. StockAnalysis shows trailing revenue growth above 400% and a 52-week range of HK\$186.20-1,330. Of the six names here, it is the furthest below its high. It shows the gap between product growth and shareholder outcome most starkly.

WATCH NEXT

Interim disclosure of gross margin on API revenue, if it provides one.

688256.CH

Cambricon Technologies

H1 net profit RMB2.3bn
7 Aug 2026

The domestic accelerator designer, and the clearest winner of China's push to rely on its own chips. First-half revenue rose 108% to RMB6bn and net profit rose 122.6%; second-quarter net profit of RMB1.3bn was up from RMB1bn in the first. It is the "landlord" of the analogy, and it is priced as one.

WATCH NEXT

Third-quarter results, and whether Nvidia inference chips from December slow its order book.

0981.HK

SMIC

HK\$63.35
late Sep 2026

The foundry that manufactures Huawei's Ascend processors. Second-quarter revenue reached \$3.01bn and gross margin improved to 25.3%, and management says demand is far exceeding earlier expectations. It is the least marked-down name in the complex, and on about 62 times trailing earnings the most richly valued.

WATCH NEXT

Capacity additions at advanced nodes and any change in US equipment controls.

The bottom line

What was repriced in September was not China's ability to build AI. That looks stronger than ever: a frontier-adjacent model at sixty cents per million output tokens, half the traffic on the largest open marketplace, and domestic chip designers doubling their profits. What was repriced was the assumption that winning the product would mean winning the profit. The market has marked down the companies closest to the models and spared the companies that sell them capacity.

The evidence is mixed, and fairly so. Platform cloud revenue is growing fast, the labs' revenues are rising at triple digits, and valuations at Tencent and Alibaba are no longer demanding. Against that, American labs have just shown they will fight on price at the bottom of the market, and the pure-play labs are still funding themselves with new shares. The reasonable inference is that China AI has become a volume story whose returns flow to bottleneck owners and to buyers, until someone shows pricing power. The next test is Alibaba's September quarter. If cloud growth holds while capex growth slows, the platforms start to look like the answer. If not, the landlord remains the only stall making a living.

Editorial note. This Market Field Guide is educational commentary prepared for internal use at Heyokha Brothers. It is not investment advice, a recommendation or an offer to buy or sell any security. Named companies are illustrative examples of how parts of a theme behave and are not recommendations. External figures are attributed and dated as of the sources given, and may have changed since. Internal Wiki views are the opinions of the named research providers, labelled as views rather than forecasts, and have not been independently verified. Where internal and external evidence conflict, the conflict is stated, not resolved.

Sources and update notes

(a) Internal. Heyokha Wiki theme page *China AI* (created 23 June 2026; last reviewed 14 September 2026), sections Thesis, The bull case, The bear case & what breaks it, Evidence & data, Contradictions, Evolution, and August 2026. Provider issues relied on: White Box (Ignacio de Gregorio), 10 June 2026 (US-China myths; capital cost per gigawatt), 18 June 2026 (GLM-5.2; 80% commodity-token view), 28 August 2026 (GLM-5.3-Flash), 14 September 2026 (DeepSeek v4.1 Flash; "top 3 lab"; ~200x price gap) and 23 September 2026 (frontier price wars; Harvey margins; Chinese accelerators unproven for frontier training); 13D Research WILTW, 16 and 30 January 2025 (original thesis) and 6 August 2026 (commoditisation as the solar and battery playbook); Greed & Fear (Christopher Wood), 27 August 2026 (usage share) and 24 September 2026 (small models versus the hyperscaler build-out); Marc Faber, 1 August 2026 (cheap-tokens strategy).

(b) Limitation on the internal view. The *China AI* page draws on 13 listed sources, and 10 of them are White Box issues by a single author. Most of its 2026 argument about cost, price and adoption therefore rests on one analyst's reading, supplemented by 13D and occasional Greed & Fear and Marc Faber material. The page says little about equity returns, valuation or the balance sheets of listed Chinese AI companies. Its

bear section concerns model capability rather than shareholder economics. Several of its data points (benchmark scores, cost-per-task figures, Harvey's margins) come second-hand through the newsletter, not from primary filings.

(c) External sources. [1] Alibaba Group, June Quarter 2026 results, 20 August 2026: revenue RMB268.95bn (+9%); cloud RMB48.44bn (+45%); AI-related product revenue RMB12.38bn; capex RMB67.68bn (+75%); adjusted EBITA RMB27.33bn (−30%); net income RMB10.44bn (−75%); free cash flow −RMB44.67bn. [2] SCMP, 12–13 August 2026: Tencent Q2 revenue RMB204.8bn (+11%), adjusted net income RMB68.4bn, capex +176%. [3] SCMP, 7 August 2026: Cambricon H1 revenue RMB6bn (+108%), net profit RMB2.3bn (+122.6%), Q2 net profit RMB1.3bn. [4] DeepSeek API change log: V4.1-Flash released 10 September 2026 with price reductions; peak/off-peak pricing effective 16 August 2026. Powerdrill summary of published V4.1-Flash pricing: input \$0.15 off-peak / \$0.30 peak; output \$0.60 / \$1.20 per million tokens. [5] SiliconANGLE, 22 September 2026: Claude Opus 5.5 at \$4/\$20, cached \$0.20 (−60%), 20% cut; GPT-6 Sol \$2/\$10 and Luna \$0.10/\$0.50, about 50% below predecessors. [6] Yahoo Finance, reporting a CNBC analysis of OpenRouter data, 8 July 2026: Chinese models 46.4% of routed tokens, US 35.7%, DeepSeek 17.6%, Qwen 13.9%; first half 2025 4.5%. [7] StockAnalysis quotes: Alibaba HK\$107.70 (28 Sep; 52-week range 88.65–186.20; forward P/E 16.0); Tencent HK\$436.60 (25 Sep; range 411–683; P/E 14.9); Baidu HK\$87.20 (18 Sep; range 85.90–161.20); Z.ai HK\$780 (18 Sep; range 116.10–2,980); MiniMax HK\$280.40 (24 Sep intraday; range 186.20–1,330); SMIC HK\$63.35 (retrieved 29 Sep, quote date not shown; range 49.32–93.50; P/E 62.2; Q2 revenue \$3.01bn, gross margin 25.3%). [8] Invezz, 14 September 2026: Z.ai −10.5% on a \$5bn raise (about \$2bn of shares at HK\$714, a 10% discount to HK\$793, plus about \$3bn of convertibles); prior \$4bn raise in July; H1 revenue RMB953.9m (+400%); Macquarie models losses into 2030; Jefferies caveats. [9] The Bamboo Works, 18 June 2026: Zhipu and MiniMax peaks of HK\$1,993 and HK\$840 on 29 May. [10] Business Today (Malaysia), 26–27 September 2026: Hang Seng Index 24,510.09 on 25 September, third weekly loss; weekly moves Alibaba −4.2%, Tencent −2.4%, MiniMax −4%, Z.ai −12.4%. [11] CSIS and SCMP, 25–26 September 2026: Trump–Xi summit 23–25 September; \$30bn of goods each way given favourable tariff treatment; AI dialogue and incident channel created; APEC in China in November. [12] Asia Times, September 2026: RTX PRO 5500 shipments from late December at about 500,000 a quarter; ByteDance and Tencent about 10,000 H200s each by August against 75,000 licensed. [13] AI News, 2026: Nvidia China revenue about 5% of total, guidance assumes zero data-centre sales to China. [14] Alizila, 22 September 2026: Apsara conference, 20GW by 2032, Zhenwu V900 at three times its predecessor's performance.

(d) Editorial choices made in this automated run. The topic was chosen by scoring every entity tagged in the 15 Wiki source summaries published 8–24 September 2026: three points within three days (none qualified), two within seven, one within 21, and three bonus points for a title mention. The higher-scoring themes, in order Gold (15), AI Capex (13), Oil & Gas (11), India Equities (10), Critical Minerals (9), Copper (8) and China Equities (6), all have editions in the archive within the past 30 days. Food Security (7) was also skipped, on the judgement that the 31 August "Grain Supply Shock" edition covered the same theme under a different name. China AI tied at six points with Iran War and Gold Miners. It was preferred because its score includes the hot signal: DeepSeek, one of the page's aliases, headlines the 14 September White Box issue. It also gains from the 23 September price-war issue, and it has no edition in the past 30 days (the nearest is China Equities on 22 September). The cooldown check read the full archive folder on the desktop through the device bridge, and filenames in both the old and current formats were parsed. No tool failures affected content. Several external pages (CNBC, one SCMP article) were blocked to the fetch tool, and figures that could only be found there were dropped.