

THE EQUITY EXPRESSION OF THE AI PRICE WAR THESIS

The Buffet That Started Weighing Plates

In the ten days to 30 September the three largest Western model makers all put a frontier-class model on the price list at \$2 per million input tokens and \$10 per million output. In the same week OpenAI halved what its \$200-a-month subscribers get, from 30 October onward. Tokens are getting cheaper; the heaviest users are about to pay more. The question is no longer whether AI prices are falling, but who ends up holding the bill when they do.

A reframe worth holding onto

The internal view, from White Box's Ignacio de Gregorio, is that the price war has two fronts that move in opposite directions: the per-token price, which keeps falling, and the shape of the contract, which is being pushed from flat subscriptions to metered consumption. In his words, subscriptions "are simply not worth it for the vendor" and "AI will end up being consumption-based." Read that way, the cuts are not a surrender of pricing power but a change in where it is exercised.

Method: external figures (model list prices, subscription terms, share prices, company financials, central-bank decisions) are taken from dated public reporting and company documentation, verified by web search on 2 October 2026 HKT and cited by number in the sources section. Internal views come from the Heyokha Research Wiki and are attributed to the named provider; they are opinions, not forecasts, and are labelled as such.

Picture an all-you-can-eat restaurant whose ingredient costs are falling every month. The owner cuts the à la carte prices, and the dining room fills. But at the buffet a handful of regulars have started arriving at noon and leaving at closing time, each eating what a hundred ordinary diners would. The flat entry ticket that once made the place easy to run now loses money on exactly the most loyal customers. So the owner does two things at once: lowers the price of every dish on the menu, and puts a scale at the buffet line, halving what the \$200 ticket allows and selling a \$500 ticket for the people who will not stop. Both moves are rational. Only one of them looks like a price war.

THE STORY IN 60 SECONDS

\$2 / \$10

Per million input / output tokens:
GPT-6.1 Sol and Claude Sonnet 5.5 (29 Sep), Gemini 4 Argon introductory (30 Sep). *Company pricing via TheVibeGate, DataCamp; as of 30 Sep 2026.*

20× → 10×

ChatGPT Pro 200 allowance as a multiple of Plus, from 30 Oct; new \$500 tier added. *The Next Web, 29 Sep 2026.*

\$518bn

Anthropic's ten-year compute obligations, ~80% non-cancellable, against \$4.6bn of 2025 revenue. *Reuters review of prospectus, via The Next Web, 29 Sep 2026.*

WHAT ACTUALLY HAPPENED

Three cuts, one squeeze, and a prospectus

The first move came on 22 September. Anthropic released Claude Opus 5.5 at \$4 per million input tokens and \$20 per million output, 20% below Opus 5, with cached input down 60% to \$0.20. The company said the model costs 40% less to run than its predecessor on typical workloads. Roughly ninety minutes later OpenAI released GPT-6 Sol at \$2 and \$10, half Opus's list price. The independent benchmarker Artificial Analysis put Opus 5.5 at a score of 51.2 for \$1.34 per task and Sol at 47.5 for \$1.06. The cheaper model was not the better one; it was the better one per dollar.

The second wave arrived a week later. At its developer day on 29 September OpenAI priced GPT-6.1 Sol at \$2 and \$10, which the company describes as a fifth of the token price of its flagship Astra, with cached input at \$0.10. Anthropic's mid-tier Sonnet 5.5 sits at the same \$2 and \$10. On 30 September Google announced Gemini 4 Argon at an introductory \$2 and \$10, rising to \$4 and \$20 after a promotional period it has not defined, with a 95% discount for cached input. Argon is for now restricted to cybersecurity partners in Google's Fairwind programme. Within a fortnight, the \$2/\$10 point stopped being a promotion and became a reference price.

The squeeze sits on the other side of the same keynote. From 30 October the ChatGPT Pro 200 plan's included usage falls from twenty times to ten times the Plus allowance, and weekly GPT-6 Pro messages drop from 200 to 100. A new \$500 tier offers 25 times Plus and an "Ultrafast" mode generating up to 300 tokens a second in Codex. Existing subscribers receive a one-time \$2,500 usage credit that expires at year-end. For a power user, matching last month's consumption now costs more than twice as much.

The third event put numbers under the argument. A Reuters review of Anthropic's draft prospectus, reported on 28 and 29 September, showed 2025 revenue up roughly twelvefold to \$4.6bn, compute and infrastructure costs of \$7.33bn, an operating loss of \$8.06bn and a net loss of \$42bn that includes a \$34bn non-cash charge tied to financing instruments. Nearly a quarter of 2025 revenue came from two customers. Revenue then reached \$4.73bn in the first quarter of 2026 and \$11.5bn in the second, and the company expects a second consecutive quarter of operating profit on an adjusted basis. The listing is expected after the November mid-term elections at a valuation reported above \$2tn.


THE ESSENTIAL DISTINCTION

A falling *price per token* is not the same as a falling *price per customer*. Agents and long reasoning chains multiply the tokens each task consumes, so the list price can halve while the bill for a heavy user doubles. The price war is real at the token; at the contract, the vendors are quietly taking pricing power back.


THE MECHANISM

Follow one task through the economics

The simplest way to see why both things can be true at once is to trace a single unit of work, one agentic task, from the price list to whoever pays for it. Each step has its own number, and the numbers do not all move the same way.



features onto an open-weight model and restored positive margins. The lesson the vendors appear to have drawn is that a flat fee is a free option written to the customer, and the most valuable customers are the ones most likely to exercise it.




Source: Reuters review of Anthropic's draft prospectus, as reported by The Next Web and Fortune, 29 Sep 2026. Ratios are this edition's arithmetic: \$7.33bn / \$4.6bn; \$518bn / (\$11.5bn × 4). The prospectus is a draft and figures may change before listing.

Those three numbers are the hinge of the whole debate. If the \$518bn is to be earned back, inference has to stay profitable while prices fall. That is possible only if efficiency gains (smaller looped models, cheaper reinforcement-learning runs, aggressive caching) cut the cost of serving a token faster than competition cuts its price. Anthropic's own claim that Opus 5.5 costs 40% less to run while its list price fell 20% is, on its face, exactly that: a price cut smaller than the cost cut. The price war, in other words, may be partly a pass-through of efficiency rather than a surrender of margin.

THE CENTRAL COMPARISON

Where the frontier price list sits now

Output tokens are the expensive half of any bill and the half agents consume most heavily. Ranked by list price per million output tokens, the frontier has compressed into a narrow band, with a long tail of much cheaper models beneath it.



Bars are illustrative, scaled to Opus 5 = 100%, and show list prices only; they ignore caching, batch discounts and tokens consumed per task, which is where most of the real cost difference lies. Opus, Sol and Sonnet prices verified externally (Implicator 22 Sep; TheVibeGate 29 Sep; DataCamp 30 Sep 2026). Grok 4.7 and GPT-6 Luna are internal Wiki figures (White Box, 23 Sep 2026) not independently re-verified.

A USEFUL MENTAL MODEL

Think of the model makers as utilities that have discovered their tariff was wrong. The per-unit charge is falling because generation is getting cheaper and rivals are building plants. But the flat-rate household plan, priced for average use, met customers who run a smelter in the garage. *The fix is not a higher unit price; it is a meter.*

WHAT HAS ACTUALLY SHIFTED

Six changes, each with two edges

WHAT HAS CHANGED	WHAT IT LOOKS LIKE IN THE DATA	WHY IT CUTS BOTH WAYS
Incumbents now cut first	Opus 5.5 down 20% on 22 Sep; GPT-6 Sol at half Opus's list price the same day.	Competition is working for buyers; but cuts paired with a 40% running-cost reduction may leave margins intact.
A reference price has formed	Three labs at \$2/\$10 for a frontier-class or near-frontier model by 30 Sep.	Convergence can signal a floor as easily as a race to zero; Argon's list price doubles after its promotion.
Heavy-user subscriptions repriced	Pro 200 allowance 20× to 10× Plus from 30 Oct; \$500 tier at 25×.	Effective price rises for the users who value AI most; it also tests whether they defect to cheaper or free rivals.
Free is now a competing price	Meta's Muse is free for most uses, with \$20 and \$100 monthly tiers; Meta shares rose 27% in September.	An advertising balance sheet can subsidise indefinitely; it also means Meta's AI spend has no direct revenue line yet.
Lab economics are now public	Anthropic: \$4.6bn 2025 revenue, \$8.06bn operating loss, \$11.5bn Q2 2026 revenue.	Growth is extraordinary; so is concentration, with two customers near a quarter of 2025 revenue.
Money got dearer	Fed raised rates 25bp to 3.75–4.00% on 16 Sep, the first hike since July 2023.	Raises the cost of carrying long compute commitments; also pressures valuations that rest on distant cash flows.

BULL AND BEAR, STATED FAIRLY

Is model-layer pricing power really eroding?

Here the "bull" case is the internal thesis: that pricing power in the model layer is eroding faster than the capex committed to serve it assumes. The "bear" case argues the thesis is overstated. The Wiki's own bear section is thin and largely about tactics, so the bear column below was built mainly from external reporting and from tensions inside the companies' own disclosures.

The price war is real and durable	Pricing power is moving, not vanishing
• The cuts are broad and fast. Every major Western lab moved within ten days, and the \$2/\$10 point recurred across three vendors by 30 September. • Pressure comes from below. The Wiki (White Box) records Chinese open-weight models and Meta's pricing forcing labs to defend share with price; Harvey restored margins by switching to an open-weight base model.	• Costs are falling faster than prices. Anthropic says Opus 5.5 runs 40% cheaper than Opus 5 while list prices fell only 20%; on that arithmetic, unit margin widened. • The contract is being repriced upward. OpenAI's Pro 200 change roughly doubles the effective price for its heaviest subscribers from 30 October.
• Buyers are learning they need less. de Gregorio's view is that "you don't need frontier tokens for most features"; cheap classifiers and small models handle routine automation. • The economics are stretched. Anthropic spent \$1.59 on compute per \$1 of 2025 revenue and carries \$518bn of obligations; price cuts make that harder to earn back, not easier.	• Revenue is accelerating, not stalling. Anthropic's quarterly revenue went from \$4.73bn to \$11.5bn between Q1 and Q2 2026, with adjusted operating profit expected for a second quarter. • Some list prices are temporary. Argon's \$2/\$10 is introductory, doubling to \$4/\$20, and the model is not yet generally available.

- **Demand may be cooling.** Greed & Fear (Chris Wood) cites a third-party index showing Anthropic server utilisation down 26% from its January peak, read as overbuilt capacity.
- **A free competitor sets the ceiling.** Meta's Muse is free for most uses, funded by advertising; no pure model vendor can match that indefinitely.

- **Holding price is still possible.** SpaceXAI kept Grok 4.7 at \$2/\$6, and frontier capability gaps on specific tasks (terminal and coding benchmarks) still differ by several points between vendors.
- **Investors are not pricing a collapse.** The leaked prospectus points to a valuation above \$2tn, up from \$965bn in May; whatever the merits, the market is not treating the layer as commoditised.

WHERE THE VIEWS PART

Internal and external evidence conflict on demand. The internal view from Greed & Fear reads the price cuts as a symptom of weakening corporate demand, citing a congestion index that puts Anthropic's server utilisation 26% below its January peak and a backlash against "tokenmaxxing". The externally reported prospectus shows the same company's revenue rising from \$4.73bn in the first quarter to \$11.5bn in the second. Both can be partly true: a queueing-based index measures waiting time, which falls if capacity is added faster than load, and the index comes from a single vantage point that Wood himself flags. But one says demand is softening and the other says it more than doubled in a quarter. This edition does not resolve it. White Box and Greed & Fear also disagree with each other: de Gregorio reads the cuts as competition working, Wood as demand faltering.

WHAT TO WATCH FROM HERE

The dashboard

SIGNAL	PLAIN-ENGLISH READING	WHY IT MATTERS
Pro 200 change takes effect, 30 Oct	Do heavy users pay up, downgrade or leave for Meta and others?	The first live test of whether metered pricing sticks for the most valuable customers.
FOMC, 27–28 Oct	A second hike after September's would raise the cost of long-dated commitments again.	Compute obligations are fixed; the discount rate applied to the revenue that must pay for them is not.
Hyperscaler Q3 results, late Oct	Capex guidance for 2027 against the ~\$990bn consensus Wood "severely doubts" (internal view).	The supply side of the price war: more capacity means more pressure to fill it cheaply.
Argon general availability	When the introductory price ends and whether \$4/\$20 holds.	Tells whether \$2/\$10 is a floor or a promotion.
Anthropic final prospectus	Gross margin by quarter, customer concentration, cancellable commitments.	The first audited look at whether inference stays profitable as prices fall.
Open-weight releases from China	Price and capability of the next DeepSeek, Kimi or GLM release.	The pressure from below that, on the internal view, drives the whole cycle.

THREE POSTURES, NOT ONE ANSWER

Own the toll roads

Express the theme through the suppliers of compute, on the logic that cheaper tokens mean more tokens and every token runs on someone's silicon. **Cost:** if Wood is right that falling prices reflect weakening demand, capex is exactly what gets cut, and the toll roads are the most exposed to a reset in 2027 spending plans.

Own the distributors

Express it through companies that can give AI away because another business pays for it: advertising, search, cloud. They benefit if models commoditise. **Cost:** they are also among the largest capex spenders, so they carry the stranded-asset risk without a model-layer revenue line to show for it.

Wait for the audited margin

Treat the model layer as unproven until a prospectus shows inference gross margins quarter by quarter, and hold the theme only through its second-order effects. **Cost:** if the metered model works and revenue keeps doubling, the waiting period is where most of the repricing happens.

The companies below are illustrative of how the theme shows up in listed markets. They are not recommendations, and nothing here is a view on whether any security should be bought, held or sold.

NVDA

Nvidia

\$228.38
close 30 Sep 2026

The supplier on whom the price war ultimately rests. Fiscal Q2 revenue rose 106% year on year, and on 30 September the board added \$150bn to its buyback, taking authorisation to \$235bn through 2028. The shares trade on about 25 times forward earnings. Cheaper tokens help volume; a capex reset would hurt orders.

WATCH NEXT

Hyperscaler 2027 capex guidance in late-October results.

META

Meta Platforms

\$725.18
close 30 Sep 2026

The price aggressor. Muse, launched 8 September, is free for most uses with \$20 and \$100 monthly tiers, and the shares rose 27% in September. Its pricing sets the consumer ceiling for every other lab. The internal view (Greed & Fear) notes Meta's talks to lease spare capacity to Anthropic as a sign of overbuild.

WATCH NEXT

Whether Muse's paid tiers convert, and any capex revision.

GOOGL

Alphabet

c.\$340.6
reported 2 Oct 2026

Gemini 4 Argon entered at an introductory \$2/\$10 with 95% off cached input. The shares rose as much as 1.7% after the announcement, then gave it back on a report of internal doubts about coding performance, down about 2.5% over 24 hours. Alphabet is also among Anthropic's largest compute counterparties, at \$111.1bn.

WATCH NEXT

Argon general availability and its post-promotion price.

AVGO

Broadcom

\$351.19
close 30 Sep 2026

The single largest line in Anthropic's commitments: \$161.2bn of equipment leases, per the prospectus. Custom accelerators are the route by which labs try to cut serving cost faster than price, the arithmetic the bear case depends on. The shares fell 1.1% on 30 September.

WATCH NEXT

Final prospectus detail on which commitments are cancellable.

CRWV

CoreWeave

c.\$86–88
intraday 30 Sep 2026

A compute landlord that, in de Gregorio's framing, avoids training costs and so reports positive operating cash flow while free cash flow stays negative (about –\$5.74bn in the latest quarter). Reported backlog is about \$104bn. The shares traded well below a September high above \$104 after a neocloud sell-off on 25 September.

WATCH NEXT

Renewal pricing on GPU contracts as token prices fall.

AMD

Advanced Micro Devices

\$611.76
close 30 Sep 2026

The alternative accelerator supplier. If labs compete on cost per token, a second source of compute gains bargaining value. The shares rose 0.7% on 30 September. The internal view (Greed & Fear) adds a caveat: small models may shift some inference toward cheaper CPUs.

WATCH NEXT

Inference share wins announced at Q3 results.

The bottom line

What was repriced in late September was the token, not the customer. List prices for frontier-class output fell by 20–50% in a fortnight and converged at \$2/\$10, while the heaviest subscription at the largest lab became roughly twice as expensive per unit of use. The vendors are conceding the rate card and reclaiming pricing power through the meter.

The evidence is mixed. Anthropic's disclosures show both stretched economics, at \$1.59 of compute per dollar of 2025 revenue and \$518bn committed, and a business whose revenue more than doubled in a quarter; the internal view reads demand as softening. The reasonable inference is that the price war is partly efficiency passed through and partly genuine margin pressure, with the split still unknown. The next test is 30 October, when heavy users decide whether to pay the meter or walk.

Editorial note. This is educational commentary prepared for internal use at Heyokha Brothers. It is not investment advice and not a recommendation to buy, hold or sell any security. Named companies are illustrative. External figures are dated and sourced; they may since have changed. Internal Wiki material reflects the opinions of named third-party research providers and is labelled as views, not facts or forecasts. Where internal and external evidence conflict, the conflict is stated, not resolved. Disclosure: this edition was drafted by an automated AI model developed by Anthropic, one of the companies discussed; it has been written to the same standard of attribution as every other company here.

Sources and update notes

Internal. Heyokha Research Wiki theme page *AI Price War* (created 25 Sep 2026; last reviewed 1 Oct 2026), sections Thesis, The bull case, The bear case & what breaks it, the dated September 29–30 update, Evidence & data and Contradictions. Concept page *AI Margins* (Harvey case; lab-side prospectus case). Provider issues relied on: White Box (Ignacio de Gregorio), "Frontier price wars, decision models, & Bonsais", 23 Sep 2026; White Box, "Gemini 4 Argon, FTC Investigations, SI Accords, & Dots", 1 Oct 2026; Greed & Fear (Chris Wood), "The virtues of small over large", 24 Sep 2026. Background: White Box, "How Agents are Shaping AI, Part 1", 28 Sep 2026.

Limitation on the internal view. The AI Price War page rests on three source issues from two providers, and two of the three are from a single author at White Box. Its thesis, its bull case and almost all of its pricing data come from that one newsletter; the only counterweight is a single Greed & Fear issue whose key datum (the 26% utilisation decline) comes from a third-party congestion index measured from one user's vantage point. None of the 13D, Solid Ground, Marc Faber, HS Dent or Gary Shilling material the Wiki carries elsewhere feeds this page. The page is ten days old. Treat it as a well-argued pair of opinions, not a consensus.

External. (1) Implicator, "Claude Opus 5.5 price cut, cyber routing and GPT-6 Sol costs", 22 Sep 2026: Opus 5.5 \$4/\$20 vs Opus 5 \$5/\$25; cache reads \$0.20 vs \$0.50; 40% lower running cost; GPT-6 Sol \$2/\$10; Artificial Analysis 51.2 at \$1.34/task vs 47.5 at \$1.06. (2) The Next Web, "OpenAI halves Pro 200 usage and launches a \$500 ChatGPT plan", 29 Sep 2026: allowance 20× to 10× Plus from 30 Oct; GPT-6 Pro messages 200 to 100 a week; \$2,500 credit; Ultrafast up to 300 tokens/s. (3) TheVibeGate, 29 Sep 2026: GPT-6.1 Sol \$2/\$10, cached \$0.10; Sonnet 5.5 \$2/\$10. (4) DataCamp, "Gemini 4 Argon", announced 30 Sep 2026: \$2/\$10 introductory, \$4/\$20 standard, 95% cache discount, Fairwind-only access. (5) The Next Web, citing Reuters, 29 Sep 2026, and Fortune, 29 Sep 2026: Anthropic 2025 revenue \$4.6bn, operating loss \$8.06bn, net loss \$42bn incl. \$34bn charge, compute costs \$7.33bn, \$518bn obligations (~80% non-cancellable; Google \$111.1bn, Broadcom \$161.2bn), two customers near 25% of revenue, Q1 2026 \$4.73bn, Q2 2026 \$11.5bn, valuation above \$2tn vs \$965bn in May, listing after November mid-terms; CNBC, 28 Sep 2026 (headline only; page not retrievable). (6) Yahoo Finance, "Stock market today, Sept. 30", 30 Sep 2026: NVDA \$228.38 (+0.51%), buyback to \$235bn, AMD \$611.76, AVGO \$351.19 (−1.10%), S&P 500 7,654. (7) Blockonomi, 1 Oct 2026: META \$725.18 close 30 Sep, +27% in September, Muse \$20/\$100 tiers. (8) ABC Money, Oct 2026: Alphabet up to \$346.52 after hours 30 Sep, c.\$340.6 on 2 Oct report, −2.51% over 24h. (9) Timothy Sykes news, 30 Sep 2026: CRWV \$86–88 intraday, ~\$104bn backlog, FCF ~\$5.74bn; 24/7 Wall St., 25 Sep 2026, neocloud sell-off. (10) Kiplinger, Sep 2026: Fed +25bp to 3.75–4.00% on 16 Sep; next FOMC 27–28 Oct. (11) Bloomberg via The Redline (X), Sep 2026: Harvey gross margin ~50% to ~50% by June, token use up twentyfold.

Editorial choices made in this automated run. The topic was chosen by scoring every Wiki entity across 14 source issues published 11 Sep–1 Oct 2026 (3/2/1 points for appearance within 3/7/21 days, plus 3 for a headline mention). Gold (16), AI Capex (14), Critical Minerals (12), Oil & Gas (11) and India Equities (9) scored higher but were all published as editions within the last 30 days, so were skipped. AI Price War and Gold Miners tied at 9; AI Price War was chosen because it appears in the newest issue (White Box, 1 Oct) and was a headline subject on 23 Sep, while gold had run as an edition on 28 Sep. The archive folder was read in full via the device bridge; nine older files use a legacy "YYYY-MM-DD - Topic" naming and were parsed manually, and none affected the cooldown. Some figures in the internal page could not be re-verified externally (Grok 4.7 pricing and token counts, GPT-6 Luna pricing, the 26% utilisation index) and are labelled as internal. The CNBC article returned an access error; its figures were taken from other outlets reporting the same Reuters review.

Prepared 2 October 2026 · Heyokha Brothers Market Field Guide · Commentary, not investment advice.