

INVESTMENT LETTER

The AI Bubble Is Real— But Not in the Way Most People Think

A genuine technological revolution, an investable scarcity premium and an emerging infrastructure risk

AS OF	ASSESSMENT	RISK SCORE
29 July 2026	Selective bubble	6.5 / 10

THE CENTRAL THESIS

AI is a genuine general-purpose technology capable of supporting a multi-trillion-dollar commercial ecosystem. Yet scarce, liquid exposure has pushed capital into a narrow group of bottleneck owners and platforms—creating a premium that is partly rational, and partly a bet that today’s scarcity will endure.

The question is not whether artificial intelligence becomes abundant. It almost certainly will. The question is: **who will still possess pricing power when it does?**

Prepared from the integrated AI Bubble article and July 2026 dashboard supplied in this conversation.

Executive summary

Bottom line. The evidence supports neither a dismissive “AI is all hype” view nor a complacent “the winners are obvious” view. AI demand and revenue are real; the risk lies in assuming that the companies building the infrastructure today will retain their scarcity economics after supply, competition and falling unit costs reshape the market.

What the data say	July reading	Investment implication
Demand and revenue	 Strong	The core commercial engine is real.
Margins and cash generation	 Healthy	Leading platforms can still self-fund much of the build-out.
Capital discipline	 Deteriorating	Capex is moving faster than cash conversion.
Market concentration	 Very high	Index and thematic exposure remain structurally fragile.
Valuation	 Uneven	Risk is concentrated in high-expectation pure plays.
Overall bubble risk	6.5 / 10	Selective bubble within a genuine supercycle.

INVESTMENT CONCLUSION

The most likely path is a real boom with a selective financial correction: the strongest platforms grow, while high-multiple infrastructure, semiconductor and application names face much more demanding economics.

Four conclusions for investors

- US\$10tn is most plausible as a broad AI-enabled ecosystem, not near-term vendor revenue.
- Scarcity matters only where it is defended by ecosystem, distribution, data or operational reliability.
- The critical metric is incremental AI revenue divided by incremental AI capital expenditure.
- Value may migrate between chips, software, infrastructure, applications and distribution.

At a glance: July 2026 dashboard

Category	Score	Signal
AI demand	9 / 10	 Strong
Revenue conversion	9 / 10	 Strong
Profit margins	8 / 10	 Healthy
Adoption breadth	6 / 10	 Developing
Valuation discipline	5 / 10	 Uneven
Capital discipline	3 / 10	 Deteriorating
Market concentration	8.5 / 10 risk	 High
Infrastructure overbuild risk	7 / 10 risk	 Rising

1. A US\$10tn market: plausible, but only with the right definition

The US\$10tn thesis is best understood as a **commercial ecosystem** rather than a single, directly measured revenue pool. The distinction matters because economic value, AI-enabled revenue and direct supplier revenue are not interchangeable.

Measure	What it captures	How to interpret US\$10tn
Direct AI revenue	Chips, servers, cloud, models, software, agents and services.	An aggressive long-term target; difficult to assess without double-counting.
AI-enabled revenue	Sales and savings generated by firms using AI in their products and operations.	Plausible as a broad ecosystem framing.
Economic value created	Productivity, lower costs, consumer surplus and new output.	Potentially larger; not the same as vendor revenue.

The practical investment implication is straightforward: **an enormous end market does not automatically validate today's valuation.** What matters is which layer captures the profit pool, whether it can protect its margins, and how much capital it takes to remain competitive.

2. The revenue is real—and that changes the comparison with 2000

The strongest defence against the “all hype” interpretation is that the leading companies are already producing revenue, margin and operating cash flow at extraordinary scale. According to the dashboard source material, Nvidia reported US\$81.6bn of quarterly revenue and approximately 75% gross margin; Microsoft cited an AI business above a US\$37bn annual revenue run rate; and Google Cloud reported US\$24.8bn of quarterly revenue, with 82% growth.

Company / period	Illustrative evidence	What it shows
Nvidia, Q1 FY27	US\$81.6bn revenue; data-centre revenue US\$75.2bn; ~75% gross margin.	AI infrastructure demand is already material.
Microsoft, Q3 FY26	US\$82.9bn quarterly revenue; AI run rate > US\$37bn.	Distribution and enterprise relationships matter.
Alphabet, Q2 2026	Google Cloud US\$24.8bn, up 82%.	Cloud demand is converting into reported revenue.
Amazon, Q1 2026	AWS US\$37.6bn; operating income US\$14.2bn.	A profitable platform can fund expansion internally.

IMPORTANT DISTINCTION

Real revenues reduce the risk of a classic revenue-free bubble. They do not eliminate capital-cycle risk: a company can be profitable today and still overinvest in capacity for tomorrow.

3. Why scarcity commands a premium

The investable universe is much smaller than the innovation universe. Thousands of firms are developing AI products, but relatively few are liquid, profitable, globally scalable and genuinely difficult to replicate. Capital seeking “pure AI” exposure has therefore been funnelled into a constrained set of listed proxies and a handful of highly valued private laboratories.

The scarce assets

- Advanced AI accelerators, high-bandwidth memory and leading-edge manufacturing capacity.
- Data-centre land, power, grid connections, cooling and high-speed networking.
- Large-scale cloud platforms, frontier-model talent and patient capital.
- Proprietary data, customer distribution and enterprise trust.

Scarcity supports a premium only when it can be defended. The companies best placed to retain it are those whose advantage survives beyond raw compute: software ecosystems, operational reliability, manufacturing access, customer distribution, data and low total cost of ownership.

4. The valuation lens: what high multiples actually require

A high price-to-earnings multiple does not merely predict “good growth”. It requires a company to grow quickly, retain a competitive position, avoid margin compression and still justify a credible terminal valuation years ahead. This is why the 50–200× range is closer to venture-capital mathematics than to normal public-market underwriting.

Purchase multiple	Earnings yield today	Earnings growth needed to support a future 20× P/E
20×	5.0%	No multiple compression assumed.
50×	2.0%	2.5× earnings growth.
100×	1.0%	5.0× earnings growth.
200×	0.5%	10.0× earnings growth.

The July dashboard qualification is important: the largest, cash-generating platforms are not uniformly trading at 50–200× earnings. The more extreme expectations are concentrated in selected software, alternative-chip and capacity-building businesses where current earnings are small, negative or dependent on a narrow customer set.

5. The key risk has moved from valuation to capital intensity

The critical AI-bubble question is no longer simply whether customers will pay. There is already substantial evidence that they will. The harder question is whether the profits generated by AI will justify the hundreds of billions of US dollars being committed to servers, data centres, networking and power.

Company	Latest capex / cash-flow evidence	Reading
---------	-----------------------------------	---------

Company	Latest capex / cash-flow evidence	Reading
Alphabet	Q2 capex of US\$44.9bn; free cash flow of –US\$5.9bn in the dashboard.	Capacity investment is exerting direct cash pressure.
Microsoft	March-quarter property and equipment spend of US\$30.9bn.	Strong operations, but rising capital intensity.
Meta	Q1 capex US\$19.8bn; 2026 outlook US\$125–145bn.	Spending trajectory is accelerating.
Amazon	Q1 cash capex US\$43.2bn; TTM FCF US\$1.2bn.	AI infrastructure is absorbing operating cash.

THE METRIC TO WATCH

Incremental AI revenue ÷ incremental AI capital expenditure. A rising ratio signals productive investment. A falling ratio over several quarters would point towards an overbuild.

Accounting earnings can obscure the transition. Infrastructure is capitalised first and expensed over time through depreciation. A company may show healthy operating income while its free cash flow weakens, then later face a rising depreciation burden even if new investment slows.

6. Falling costs and early adoption: both bullish and bearish

Falling inference costs are a benefit to users and a double-edged sword for suppliers. If demand expands faster than prices fall, total spending and adoption can compound. If models and compute become interchangeable faster than usage grows, society receives more intelligence for less money while supplier margins decline.

Force	Bullish reading	Bearish reading
Falling unit costs	Lower cost unlocks many more use cases.	Pricing power compresses before volume compensates.
Open models and custom chips	Broader ecosystem accelerates adoption.	Incumbent bottlenecks become less durable.
Early enterprise adoption	Substantial runway remains.	Capacity is being committed before full ROI is proven.
Physical constraints	Power and data centres support near-term scarcity.	Long contracts may be signed at peak-cycle economics.

THE UNDERWRITING CHALLENGE

The bull case requires paid usage to grow faster than unit prices fall. The bear case is not “AI fails”; it is “AI succeeds so broadly and cheaply that profits migrate away from today’s suppliers.”

7. History: four lenses for the current cycle

Analogy	What it teaches	AI relevance
Railways	Transformative infrastructure can be overbuilt and poorly financed.	Data-centre capacity may become abundant before returns are earned.
Electricity	Installation comes before organisational redesign and broad productivity.	AI benefits may arrive later than financial markets expect.
Internet & telecom	The technology thesis can be right while securities are priced wrongly.	Profitable leaders are stronger than in 2000, but capex risk remains.
Automobiles	A huge market can still produce many failed companies and shifting leadership.	Thousands of applications may become features rather than businesses.

The shared lesson is not that AI is doomed. It is that **a technology can be socially indispensable while the capital that financed its early build-out earns poor returns.**

8. A selective bubble, not a uniform one

Market layer	Current character	Primary risk
Leading profitable platforms	Expensive secular growth with genuine cash generation.	Capital intensity and slowing incremental returns.
Dominant chip suppliers	Scarcity economics and high margins.	Competition, customer concentration and cyclicity.
Cloud / data-centre infrastructure	Strong demand, physical bottlenecks.	Overbuilding, depreciation and financing terms.
Frontier-model laboratories	Strategic optionality and rapid innovation.	Cash consumption and uncertain steady-state margins.
AI applications	Fast experimentation and product proliferation.	Commoditisation, weak distribution and feature absorption.
AI-labelled, unprofitable names	Narrative-driven exposure.	Financing risk and valuation reset.

9. Scenario framework

The probabilities below are judgemental, not statistically precise. They are designed to clarify the conditions that would validate—or undermine—the thesis.

Scenario	Probability	What happens
Bull: demand absorbs the build-out	25%	AI revenue compounds faster than capex; utilisation stays high; free cash flow recovers; leading companies grow into elevated valuations.
Base: real boom, selective correction	50%	AI remains transformative, but returns become uneven; some capacity is overbuilt and high-multiple names correct sharply.
Bear: infrastructure and financing unwind	25%	Prices fall faster than paid usage grows; utilisation weakens; depreciation rises; external financing becomes scarce.

BASE CASE

The most likely outcome is not a collapse of AI demand. It is a redistribution of returns: durable platforms survive, some infrastructure economics weaken, and the market's current favourites do not all remain the end-state winners.

10. What to monitor each month**Evidence the cycle remains healthy**

- AI-related revenue rises at least as quickly as capital expenditure and depreciation.
- Free cash flow stabilises as new capacity is absorbed.
- Enterprise users report measurable returns, not just experimentation.
- Utilisation remains high and demand broadens beyond hyperscalers and model laboratories.
- Market gains broaden across applications and AI-enabled beneficiaries.

Signals that bubble risk is becoming systemic

- Capex expands much faster than AI revenue while free cash flow deteriorates.
- Data-centre utilisation or compute pricing weakens after a surge in supply.
- Depreciation rises rapidly without corresponding revenue and margin expansion.
- Debt, vendor financing, customer guarantees or off-balance-sheet vehicles become essential to keep capacity growing.
- Valuations rise while earnings estimates stagnate and market leadership narrows further.

Refined conclusion

AI is unlikely to be a fake boom. It is a real general-purpose technology that may support a multi-trillion-dollar ecosystem across semiconductors, infrastructure, models, software, robotics and AI-enabled industries. But the market is not pricing the average outcome. It is pricing a more demanding one: that today's bottleneck owners retain high market share, high margins and capital discipline as the technology scales.

That is where the scarcity thesis is strongest—and where it must be handled most carefully.

Scarcity explains the premium; it does not guarantee its permanence. As capital flows in, supply responds, technology improves and new competitors emerge. The revenue pool may continue growing even as the scarcity premium contracts.

FINAL VIEW

AI is a genuine supercycle containing pockets of bubble-like pricing. The enduring winners will be the firms that own the customer, control a defensible bottleneck, sustain returns on capital and retain pricing power after intelligence becomes abundant.