

The Information Daily Briefing — July 22, 2026

Prepared for SC · Source: The Information (Gmail newsletters + theinformation.com, logged-in) · Coverage window: last 24 hours (stories dated July 21, 2026 PDT)

Top Stories (What · Why · When · Which · How)

1. Mercor's fast growth leans heavily on a handful of AI labs

- **What:** Mercor, a three-year-old AI data-labeling startup, generated **\$614M in gross revenue in H1 2026 — up 70% versus all of 2025** — with about **91% coming from AI foundation-model makers**. It is discussing a raise that could value it at **~\$20B, double last fall's mark**.
- **Why:** The customer concentration crystallizes a rising investor worry — startups whose growth depends on a few large customers who could switch vendors or bring the work in-house. Gross margins are thin (27% in 2025, 33% in Q2 2026) because Mercor pays ~two-thirds of revenue to its contractor army.
- **When:** Reported July 21, 2026; documents cover H1 2026 with projections to 2030 (targeting ~\$69B gross revenue, 45% operating margin).
- **Which (named):** Mercor; customers OpenAI, Anthropic, Google DeepMind, Reflection AI, Thinking Machines Lab, plus Ramp, Blue Owl, Citi; rivals Scale AI, Handshake AI, Surge AI; recent acquisition Deeptune; Meta (owns 49% of Scale, paused Mercor after a March breach); margin comps GitLab, Palantir, OpenAI, Anthropic.
- **How:** Based on financial documents seen by The Information.
- **Link:** <https://www.theinformation.com/articles/mercors-fast-growth-relies-biggest-ai-companies-documents-show>

2. Meta building an OpenRouter rival ("Switchboard") to cut AI costs

- **What:** Meta's internal incubator, **AAI Labs**, is developing a model router called **Switchboard** that scores task difficulty and sends simpler requests to cheaper models — one of ~200 approved early-stage projects.
- **Why:** Inference cost is "what limits how widely we can run agents." Meta pays "top-model prices for every coding request, including the easy ones," and wants to cut billions in AI spend while turning its enormous AI investment into new products/revenue beyond ads. Capex could hit **\$145B in 2026, more than double 2025**.
- **When:** Reported July 21, 2026; per a July internal memo (team formed March 2026).
- **Which (named):** Meta; OpenRouter (the model it emulates, valued \$1.3B and fielding takeover interest); OpenAI (GPT-5 router precedent); Databricks and Palantir (own routers).
- **How:** Based on internal documents reviewed by The Information.

- **Link:** <https://www.theinformation.com/articles/metas-ai-incubator-developing-openrouter-rival-cut-coding-costs>

3. Anthropic eyes robotics; Nvidia ramps “Vera Rubin”

- **What:** Anthropic held **early (ultimately abandoned) acquisition talks this spring with Physical Intelligence** (robotics software, last valued ~\$11B), signaling robotics as a future frontier. Separately, **Nvidia’s next-gen Vera Rubin racks** are shipping as test units, priced at **\$7M–\$8M each vs ~\$5M for Grace Blackwell**, with capacity to eventually build up to 1,000 racks/day.
- **Why:** Anthropic believes “embodied intelligence” follows recursive self-improvement; robots could even help build/run data centers. For Nvidia, Vera Rubin is the next revenue engine, and it is hedging against rival AI chips by selling non-GPU components too.
- **When:** Reported July 21, 2026; Vera Rubin expected at scale in 2027.
- **Which (named):** Anthropic, Physical Intelligence (backers OpenAI, Thrive Capital); Nvidia; early rack recipients CoreWeave, Microsoft, OpenAI, Anthropic, SpaceXAI; Google (competing inference chip). Deal flow noted: Tempus AI acquiring Personalis (\$1.5B); Neo (\$100M, a16z/Bessemer); Zhipu AI’s all-domestic-chip data center; Alphabet +3% on The Information’s Google-chip report.
- **How:** Reporting/sources plus a visit to Nvidia’s Santa Clara HQ.
- **Link:** <https://www.theinformation.com/newsletters/ai-agenda/anthropics-robot-ambition-nvidia-ramps-vera-rubin>

4. New data: Anthropic lags on customer-service AI (speed & cost)

- **What:** Data from software vendor **NiCE** shows Anthropic’s models are slower (“time to first token”) and far pricier for voice/chat customer service. **Opus 4.7 took >4 seconds to first token vs ~1 second for Google’s Gemini 3 Flash, and costs 18.5x more**; Anthropic’s fast Opus 4.6 can cost ~100x the Google model. Anthropic’s edge: **accuracy (roughly half the error rate)**.
- **Why:** In voice service, milliseconds matter (a bot can pause ~2.5s max before sounding unnatural). SaaS “orchestrators” increasingly position as model-agnostic — Atlassian openly cautions against routing all data/contacts to one lab.
- **When:** Reported July 21, 2026.
- **Which (named):** Anthropic; Google (Gemini 3 Flash); OpenAI (5.4-mini, 4.1-mini, new bidirectional audio model); NiCE (clients Toyota, Allianz, Nestlé, Frontier Airlines); Atlassian (Jira agent), Microsoft (GitHub Copilot, Teams).
- **How:** Analysis of NiCE customer data.
- **Link:** <https://www.theinformation.com/newsletters/applied-ai/anthropic-best-powering-customer-service-new-data-show>

5. Venture’s AI fever goes nuclear — Aalo Atomics raising

- **What:** **Aalo Atomics** (small fission reactors sited near data centers) is courting investors for a new round — the latest in a nuclear-financing spree alongside **Valar Atomics** (raising ~\$1B at ~\$5B pre-money) and **Helion Energy** (raised \$465M at \$15B pre, fusion).

- **Why:** AI data centers need novel clean power, and unlike model makers, the AI-energy winners haven't emerged yet. Momentum followed several startups hitting a **July U.S. DOE self-sustaining-reaction (criticality) milestone**. Big VCs are now leaning in.
- **When:** Reported July 21, 2026; Aalo's \$100M Series B closed last August; DOE milestone met in July.
- **Which (named):** Aalo Atomics, Valar Atomics, Helion Energy, Antares Nuclear; Oklo (public, -40% YTD but 4x its IPO price); investors Thrive Capital, Sequoia (Shaun Maguire), Valor Equity Partners; adjacent SpaceX, OpenAI (Altman-backed nuclear ties).
- **How:** Reporting from people familiar with the fundraising.
- **Link:** <https://www.theinformation.com/newsletters/dealmaker/ventures-ai-fever-goes-nuclear>

6. The Briefing: Dorsey's "Buzz" takes on Slack; TSMC eyes 10% price hike

- **What:** Jack Dorsey's **Block** unveiled **Buzz**, a "collaboration workspace" (threads, DMs, voice, code repos, workflows) built for humans *and* AI agents, explicitly aiming to cut "dependency on Slack and GitHub." Separately, **TSMC is reportedly set to raise chip prices up to 10% next year** (Nikkei), following ASML's own ~10% equipment price hike.
- **Why:** Software isn't dead — but Buzz faces entrenched incumbents. More importantly, chip-cost inflation from TSMC/ASML flows straight to Nvidia, Apple and Google (all starved for TSMC allocation, with little room to push back), compounding data-center power and memory-cost pressures.
- **When:** July 21, 2026 (Buzz launch; TSMC/Nikkei report Tuesday).
- **Which (named):** Block, Salesforce (Slack/Benioff), Microsoft (GitHub); TSMC, ASML, Nvidia, Apple, Google, Oracle (data-center costs). In Other News: Microsoft buying Nvidia GPUs to share with Mistral AI; Google's new "Flash" Gemini models; Tesla Robotaxi expands to Tampa & Orlando; Ant International raised ~\$1.2B; Treasury Secretary Bessent backs open-source AI but threatens sanctions on IP theft; Supermicro +17% after-hours.
- **How:** Company announcements, Nikkei reporting, and earnings-call commentary.
- **Link:** <https://www.theinformation.com/newsletters/the-briefing/dorsey-wants-buzz-loosen-slacks-hold-workers>

Overall Summary — Trend, Development & Sentiment

The last 24 hours read as one continuous story about **the economics of the AI buildout**. Three developments dominate:

1. Cost inflation is the new headline. TSMC (+10%) and ASML (+10%) price hikes, Oracle's multibillion-dollar data-center cost surprises, Nvidia's pricier Vera Rubin racks (\$7–8M vs \$5M), and elevated memory prices all point the same direction — the inputs to AI are getting more expensive, not cheaper. That, in turn, is driving an industry-wide

scramble to **cut inference cost through model routing** (Meta's Switchboard, OpenRouter's takeover interest, OpenAI's GPT-5 router, Databricks/Palantir routers).

2. Capital keeps flooding AI-adjacent picks-and-shovels. Money is chasing the enablers of AI more than the model makers themselves — data labeling (Mercor at ~\$20B), nuclear power for data centers (Aalo, Valar, Helion), robotics (Physical Intelligence ~\$11B), and record U.S. equity issuance (Alphabet, SpaceX). Venture appetite that once avoided long-horizon “hard tech” like nuclear is now leaning in after real technical milestones.

3. Model differentiation and commoditization are colliding. The NiCE data is a concrete example: Anthropic wins on accuracy but loses badly on speed and price for high-volume voice work, where Google's Gemini 3 Flash and OpenAI's mini models dominate on cost. Routing-to-cheapest and open-source diffusion (Kimi K3, Zhipu's domestic-chip data center) both chip away at premium, proprietary pricing power.

Sentiment: constructively frothy. The tone is bullish on the AI buildout and its financing (“AI fever”), but layered with visible **cost discipline and margin scrutiny** — investors are now asking who *pays* for the buildout and whether the unit economics hold. Geopolitics (U.S.–China, open-source, Bessent's sanctions threat) is a persistent undercurrent rather than the day's main event.

Important Questions & Risks Being Raised

- **Customer concentration:** Is Mercor's growth durable if OpenAI, Anthropic or Google insource labeling or switch to rivals (Scale, Surge, Handshake) — or to synthetic “RL environments” that need fewer humans?
 - **Margin sustainability:** Can data-labeling (27–33% gross margin) and even the model makers hit their aggressive margin forecasts while inference and chip costs climb?
 - **Pricing power / commoditization:** If routers send “easy” work to cheap models, does premium pricing (Anthropic's) erode? Meta building its own router internally is a warning sign for third-party providers.
 - **Cost inflation pass-through:** With TSMC and ASML raising prices into a captive customer base, how much margin compression lands on Nvidia, Apple, Google and the hyperscalers?
 - **China & open source:** Do freely diffusing Chinese models (Kimi K3) and domestic-chip data centers (Zhipu) undercut U.S. control-based AI policy and export leverage? Bessent floated sanctions for IP theft.
 - **Nuclear timing/returns:** Can fission/fusion startups convert milestones into revenue on a venture timeline? Oklo's -40% YTD is a caution flag.
 - **Data governance:** Atlassian's pushback (“I don't want all my company's information going to Anthropic”) spotlights whether labs could use customer data to build competing products.
-

Potential Investment Implications (companies directly mentioned)

Informational only — not investment advice. Figures/claims are as reported by The Information.

Public equities

- **Nvidia (NVDA):** Vera Rubin's higher ASPs (\$7–8M/rack) and huge unit potential are a tailwind, but watch TSMC price hikes (input cost) and intensifying competition (Google's new inference chip). Nvidia's move to sell non-GPU components is a deliberate hedge.
- **Alphabet / Google (GOOGL):** The day's relative winner — Gemini 3 Flash is beating Anthropic on cost/speed in customer service; the "Frozen" chip could be 6–10x more efficient (stock jumped ~3%); new Flash Gemini models shipped. Reports earnings this week. Offsetting: it's a TSMC customer facing higher prices.
- **TSMC (TSM):** Pricing power (+10% reported) and "very strong" multi-year demand — a positive read for TSMC itself.
- **ASML:** Raising prices ~10% into overwhelming demand — supportive.
- **Microsoft (MSFT):** Buying Nvidia GPUs to share with Mistral AI; GitHub is flagged as "vulnerable" under AI-usage strain (a competitive risk as Block's Buzz and others circle).
- **Salesforce (CRM):** Slack faces a new entrant in Block's Buzz atop existing investor worries about AI disrupting enterprise SaaS — a sentiment headwind.
- **Block (XYZ):** Buzz adds optionality but enters an entrenched, crowded collaboration market; execution risk is high.
- **Meta (META):** Switchboard = cost discipline on a \$145B capex base; monetization ambitions are broadening. Also exposed via its 49% Scale AI stake (which paused work with Mercor rival dynamics).
- **Oracle (ORCL):** Data-center cost surprises are a margin risk to watch.
- **Palantir (PLTR):** Cited as a best-in-class margin benchmark and a router-tool developer.
- **Tesla (TSLA):** Robotaxi expansion to Tampa & Orlando continues the driverless rollout.
- **Tempus AI (TEM) / Personalis (PSNL):** M&A event — Tempus to acquire Personalis at \$1.5B.
- **Supermicro (SMCI):** +17% after-hours on higher-than-expected gross margin, though revenue near the low end (~90% YoY growth implied).
- **Oklo (OKLO):** Down ~40% YTD (still 4x IPO) — a sentiment proxy for listed nuclear.

Private / pre-IPO to watch: Anthropic, OpenAI, Mercor (~\$20B), Physical Intelligence (~\$11B), OpenRouter (\$1.3B, takeover interest), Aalo Atomics, Valar Atomics (~\$5B), Helion Energy (\$15B), Scale AI, Handshake AI, Surge AI, CoreWeave (public: CRWV), Mistral AI, Databricks, Ant Group / Ant International, Zhipu AI.

Net lean: Constructive on AI infrastructure enablers (compute, chip-equipment, power), with a bias toward **cost-efficiency winners** (Google/Gemini Flash economics, routing) over premium-priced model exposure. Rising input costs (TSMC/ASML) and commoditization argue for selectivity, not blanket AI beta.

Sources: The Information — The Briefing, Dealmaker, AI Agenda, Applied AI, and Exclusive reports (July 21, 2026), accessed via Gmail newsletters and theinformation.com. Links provided per story above.