

The Information Daily Briefing — 2026-08-14

Coverage window: last 24 hours (The Information, 12–13 Aug 2026 US time)

Friday, 14 August 2026 (GMT+8)

The Information Daily Briefing — 14 August 2026

Prepared from The Information’s newsletters and briefings received in the last 24 hours (The Briefing, The Information AM, Applied AI, AI Agenda, and Big Read features). Source articles are linked at the end of each item; some feature links sit behind The Information’s subscriber paywall.

Top Stories — What / Why / When / Which / How

1. The AI compute market is “frothy” — neoclouds ride record pricing

- **What:** Neocloud operators CoreWeave and Nebius told analysts that AI compute pricing is hitting new highs. Nebius held its first-ever capacity auction, with Nvidia Blackwell chips clearing ~15% above its previous top price; CoreWeave raised prices 25% in July and said near-term capacity is “effectively sold out.”
- **Why:** Insatiable demand for Nvidia AI chips has outstripped available data-center capacity, so customers with immediate, time-bounded needs pay steep premiums on short-term deals.
- **When:** Earnings calls 11–12 Aug; reported in The Briefing (13 Aug).
- **Which:** Nebius (NBIS), CoreWeave (CRWV), Nvidia (NVDA), SpaceX (private), with analyst commentary from MoffettNathanson and BNP Paribas.
- **How:** Auctions, short-duration high-priced contracts, and selling capacity closer to delivery to capture spot-like pricing. CoreWeave rose 19% and Nebius soared 34% on the week despite heavy cash burn.
- **Link:** <https://www.theinformation.com/briefings/nebius-shares-soar-q2-revenue-surge>

2. Nebius Q2 revenue up 454% — but cash burn balloons

- **What:** Nebius reported Q2 revenue of \$582M, up 454% year-on-year, while cash burn jumped to \$3.4B (from \$678M a year earlier) and capex hit \$5.66B. Shares jumped 17% on Wednesday.
- **Why:** Surging AI-computing demand is letting neoclouds scale revenue fast, but the land-grab for chips, servers and data centers is torching cash.
- **When:** Q2 results, 13 Aug.
- **Which:** Nebius (NBIS); CEO Arkady Volozh said Nebius could “sell today our entire 2027 capacity...if we wanted.”
- **How:** Aggressive capacity build-out funded by capex, monetized through premium short-term pricing.
- **Link:** <https://www.theinformation.com/briefings/nebius-shares-soar-q2-revenue-surge>

3. The compute crunch is squeezing “neolabs” hardest

- **What:** Startups training their own models can’t get GPUs on reasonable terms. One-year Nvidia H100 rental prices rose ~50% to \$2.60/hr (mid-June) from \$1.73 (mid-Dec); contract terms have stretched to 3–5 years with 12–18 month waits for the newest clusters, and on-demand capacity has largely “sold out.”
- **Why:** Cloud providers now push multi-year commitments to lock in revenue — untenable for a one-year-old startup being asked to sign a five-year deal worth 3x its total funding.
- **When:** Sourcing experiences from winter 2025/26 through mid-2026; AI Agenda (13 Aug).
- **Which:** Generalist (Evan Morikawa), Reflection AI (renting SpaceX capacity at ~\$150M/month plus a \$1B+ Nebius deal), Prime Intellect, Runpod, CoreWeave, Nebius, AWS.
- **How:** Neolabs scramble across 15–20 providers, sign multi-million-dollar deals, and increasingly source from secondary/offshore data centers (e.g., India via Prime Intellect).
- **Link:** <https://www.theinformation.com/> (AI Agenda, “Why the AI Compute Crunch Is Hitting Neolabs Especially Hard”)

4. Study challenges the “Chinese AI is cheaper” narrative

- **What:** An AlphaSense study found OpenAI’s GPT-5.6 Sol and Anthropic’s Opus 4.8 delivered better answers at *lower effective cost* than China’s Kimi K3 (Moonshot) and GLM-5.2 (Z.ai) on 246 financial-analysis questions — despite higher per-token prices — because frontier models use fewer tokens per task.
- **Why:** Per-token price is misleading; greater model intelligence can lower total cost per completed task, countering fears that cheap open-weight Chinese models will erode US labs’ pricing power.
- **When:** Study released early August; Applied AI (13 Aug).
- **Which:** OpenAI, Anthropic, Moonshot (Kimi K3), Z.ai (GLM-5.2), AlphaSense, model-router OpenRouter; tested via Amazon, Google, Microsoft and Cerebras.
- **How:** GPT-5.6 Sol was ~13% cheaper and ~20% higher quality than Kimi K3; Opus 4.8 ~half the cost and ~13% higher quality. AlphaSense advocates a mixed-model “router” approach for best cost/quality.
- **Link:** <https://www.theinformation.com/> (Applied AI, “When Anthropic and OpenAI Are Cheaper to Use Than Chinese Models”)

5. DeepSeek’s V4-Pro flagship lands to mixed reviews

- **What:** DeepSeek released its flagship V4-Pro model; a Vals AI leaderboard ranked it #2 among open-source models (behind Kimi K3), but testers flagged lost reasoning continuity and weak image performance — in some coding tasks it trailed DeepSeek’s own smaller V4-Flash.
- **Why:** Living up to expectations is getting harder; V4-Flash’s earlier low-cost success raised the bar.
- **When:** Official release Wednesday; The Information AM (13 Aug).
- **Which:** DeepSeek, Moonshot (Kimi K3), Anthropic (Claude Opus 5 as a price benchmark).
- **How:** DeepSeek priced V4-Pro remarkably low (\$0.435 input / \$0.87 output per 1M tokens vs Kimi K3 at \$3/\$15 and Opus 5 at \$5/\$25) even after signaling a “significant” future price hike.
- **Link:** <https://www.theinformation.com/briefings/deepseek-releases-flagship-v4-pro-model-challenge-kimi-k3>

6. AI's data land-grab: old Slack threads and IT tickets are now assets

- **What:** After AI-agent startup Warmly agreed to be acquired by HubSpot, it fielded four approaches — including from Mercor — offering up to \$300,000 to buy or license its codebase, Slack messages, GitHub/Asana content, Google Drive docs and years of staff-meeting transcripts.
- **Why:** Labs (OpenAI, Anthropic, Google) increasingly need real enterprise workflow data to train agentic models; a startup's exhaust data suddenly has resale value.
- **When:** Late June onward; Exclusive (13 Aug).
- **Which:** Mercor, Warmly (HubSpot), OpenAI, Anthropic, Google.
- **How:** Data brokers approach acquired companies for datasets the acquisition excludes, converting dormant internal records into training assets.
- **Link:** <https://www.theinformation.com/> ("Startups Find Old Slack Threads, IT Tickets Are Suddenly in High Demand")

7. Kalshi in talks to raise at \$40B valuation

- **What:** Prediction market Kalshi is in advanced talks to raise $\geq$ \$750M at a \$40B valuation, with Sequoia and Wellington in talks to co-lead. Annualized revenue topped \$4B in July (roughly doubled in ~two months) on World Cup wagers.
- **Why:** Prediction markets are booming; Kalshi has overtaken Polymarket, and early IPO conversations point to a possible listing as soon as next year.
- **When:** Exclusive; The Information AM (13 Aug). Prior round: \$1B at \$22B (May).
- **Which:** Kalshi, Sequoia (existing investor; partner Alfred Lin on board), Wellington (new), Polymarket, CFTC.
- **How:** Sports contracts drive >80% of volume; risks include heavy marketing spend, new state taxes and legal battles with state regulators alleging illegal gambling.
- **Link:** <https://www.theinformation.com/briefings/exclusive-sequoia-wellington-talks-back-kalshi-40-billion-valuation>

8. Cerebras shares fall 16% despite fast growth

- **What:** Cerebras Q2 revenue rose 74% to \$180M, but hardware sales fell to \$54M (about half the prior quarter); net loss was \$451M (incl. ~\$400M stock-based comp). Shares dropped 16% after hours; FY revenue guide raised to \$880-890M.
- **Why:** Investors focused on declining hardware revenue, customer concentration and competition rather than headline growth.
- **When:** Q2 results, 13 Aug (second earnings call since IPO).
- **Which:** Cerebras, OpenAI and Mohamed bin Zayed University of AI (each ~1/3 of revenue), plus AMD and Amazon partnerships.
- **How:** Cerebras is pushing "disaggregated inference" (running its chips alongside AMD/Amazon silicon); RPO of \$25.4B is roughly flat, with a "significant amount" tied to OpenAI. Stock is >40% below peak but above its May IPO price.
- **Link:** <https://www.theinformation.com/briefings/cerebras-shares-fall-16-decline-hardware-revenue>

9. Cisco's AI-networking boom accelerates revenue

- **What:** Cisco reported July-quarter revenue of \$17.3B, up 18%, driven by cloud providers buying AI networking chips and switches; it booked \$4B in AI product orders in Q4 alone and expects AI revenue to rise to \$7.5B this fiscal year (from \$4B).
- **Why:** Cloud providers need fast networking to maximize utilization of scarce, expensive GPUs.
- **When:** Q4/FY results, 13 Aug; guiding to 21% growth next quarter.

- **Which:** Cisco (CSCO), Nvidia (as GPU supplier), hyperscaler customers.
- **How:** Despite the beat and acceleration (from 8% growth the prior quarter), shares fell 5% on elevated expectations.
- **Link:** <https://www.theinformation.com/briefings/cisco-systems-says-cloud-providers-buying-ai-chips-switches>

10. Memory-cost inflation reaches consumers: Google lifts Pixel prices

- **What:** Google raised prices across its new Pixel 11 lineup by \$100 (Pixel 11 Pro now \$1,099; base \$899; foldable \$1,899), citing soaring memory-chip costs.
- **Why:** AI-driven demand for memory is inflating component costs across consumer hardware; Apple has already raised MacBook/iPad prices and iPhones are expected to follow.
- **When:** Pixel 11 unveiled Wednesday; The Briefing / AM (13 Aug).
- **Which:** Google/Alphabet (GOOGL), Apple (AAPL), and (by implication) memory suppliers Samsung, SK Hynix and Micron.
- **How:** Pass-through of higher input costs to consumers — a broadening signal of an AI-led memory up-cycle.
- **Link:** <https://www.theinformation.com/briefings/google-raises-prices-latest-pixel-phones>

11. Tencent's capex nearly triples on AI compute

- **What:** Tencent's Q2 capex nearly tripled year-on-year to 52.8B yuan (\$7.8B) to expand computing for training and to power its coding/productivity agents; revenue rose 11% but net profit barely grew on higher R&D.
- **Why:** Tencent is prioritizing compute for its own models (HY3 now; bigger HY4 coming) over renting capacity externally, betting on long-term gains over near-term cloud profit.
- **When:** Q2 results, 13 Aug.
- **Which:** Tencent (TCEHY / 0700.HK), plus Chinese rivals Moonshot and Alibaba.
- **How:** In-house compute allocation to WorkBuddy (agent) and CodeBuddy (coding), aiming for "state-of-the-art" with HY4.
- **Link:** <https://www.theinformation.com/briefings/tencents-capex-nearly-triples-compute-ai-models-tools>

12. Also notable

- **Kushner & Iger to buy the LA Lakers:** Thrive Capital's Josh Kushner and ex-Disney CEO Bob Iger agreed to buy control of the Lakers at a \$12.5B valuation (up from \$10B last year), pending NBA approval. (*ESPN, via The Information AM*) — https://www.espn.com/nba/story/_/id/49590362/josh-kushner-bob-iger-buy-lakers-12b
- **Big Read — "Anthropic's First Lady":** A profile of Camilla Clark, wife of CEO Dario Amodei, cast as his behind-the-scenes "emotional ballast" and one of the most consequential people in Anthropic's orbit without a formal role. — [https://www.theinformation.com/\(Big+Read\)](https://www.theinformation.com/(Big+Read))
- **Deals & debuts:** Thrive Holdings raised \$2B (D1 Capital, Altimeter, SoftBank); Lovable raised \$400M at a \$13.3B valuation (Menlo Ventures); Legora in talks at ≥\$10B; Anthropic in talks to buy world-model developer Decart AI for ~\$6B (Bloomberg); CodeRabbit \$143M at \$1.5B; Skan AI \$63M; Mindgard \$30M.

Market Trend, Development & Sentiment

The through-line is compute scarcity and pricing power. Nearly every story this cycle traces back to one dynamic: demand for Nvidia-class AI compute is running far ahead of available capacity, and that scarcity is now visibly rippling outward — into neocloud revenue and share prices (Nebius, CoreWeave), into networking demand (Cisco), into startup survival economics (neolabs), and even into consumer device prices (Pixel, and soon iPhone) via memory-cost inflation. The symbolic marker: Nvidia's market cap is now ~\$1 trillion above Apple's, having reversed a lead Apple briefly held two weeks ago.

Sentiment is bullish but increasingly two-tiered. Investors are rewarding demonstrated pricing power and revenue acceleration (Nebius +34% on the week, CoreWeave +19%, SpaceX's implied value recovering) while punishing any blemish in otherwise strong prints (Cerebras -16%, Cisco -5% despite an 18% revenue beat). The message: the bar is high, and the market is discriminating between "priced-in growth" and "genuine upside surprise."

A quieter counter-narrative is forming around efficiency and value, not just price. The AlphaSense study reframes the US-vs-China model debate away from headline per-token pricing toward total cost per completed task — a win for frontier US labs (OpenAI, Anthropic) and for model-routing tools (OpenRouter). Meanwhile DeepSeek's mixed V4-Pro reception suggests the open-weight cost advantage is not automatic, and the emerging market for enterprise training data (Mercor) shows the frontier is shifting from raw model scale toward proprietary data and real-world agentic capability.

Key Questions & Risks

- **Is compute pricing a durable moat or a cyclical spike?** The Briefing itself cautions against over-reading record prices: premiums are concentrated in *short-term* deals, while multi-year contracts (2027-2029 delivery) don't carry the same markup because hundreds of billions in capacity is being built. A capacity glut in 2027-28 is the central risk to the neocloud bull case.
- **Cash burn vs. revenue growth:** Nebius (\$3.4B burn) and CoreWeave are financing explosive growth with explosive spending. If pricing normalizes before capacity is absorbed, the leverage cuts both ways.
- **Customer concentration:** Cerebras derives ~two-thirds of revenue from just two customers (OpenAI, MBZUAI); much of its \$25.4B backlog is OpenAI-linked. Any change in a single relationship is material.
- **Neolab viability:** Multi-year, multi-hundred-million-dollar compute commitments may be structurally incompatible with early-stage model startups — a risk to the diversity of the model ecosystem and to venture returns in the segment.
- **Regulatory overhang on Kalshi:** State-level "illegal gambling" challenges and potential new state taxes could dent a business where >80% of volume is sports contracts, even as it courts a \$40B valuation and an IPO.
- **Consumer demand elasticity:** Memory-driven price hikes on phones/PCs could pressure unit volumes into a softer consumer backdrop, a downside risk for Apple and Google hardware.
- **Macro noise:** The Briefing attributes part of the Nvidia/Apple divergence to broad market dynamics (e.g., geopolitics/"the state of the Iran war") — a reminder that AI-thematic moves are riding on a volatile macro tape.

Potential Investment Implications

(For information only — not investment advice. Companies below are drawn directly from the articles.)

- **Nvidia (NVDA):** The clearest structural beneficiary — chip scarcity is translating into pricing power for it and its whole downstream ecosystem. Watch for signs the market is over-extrapolating short-term rental pricing.
- **Nebius (NBIS) / CoreWeave (CRWV):** High-torque plays on compute pricing with genuine revenue acceleration, but heavy cash burn and sensitivity to 2027+ capacity normalization. Position sizing should reflect volatility; The Information explicitly flags taking a “longer-term view.”
- **Cisco (CSCO):** A comparatively lower-beta way to play AI-networking demand; accelerating growth and a \$7.5B AI revenue trajectory, though the −5% reaction shows expectations are already elevated.
- **Cerebras:** High growth but hardware softness, ~\$400M SBC and two-customer concentration make it a higher-risk, headline-sensitive name; ~40% below peak.
- **Apple (AAPL):** Losing relative momentum to AI-levered peers and facing memory-cost margin pressure; the price-hike strategy tests consumer elasticity.
- **Alphabet / Google (GOOGL):** Dual exposure — beneficiary via AI/cloud demand, but Pixel price hikes signal input-cost pressure in hardware.
- **Tencent (TCEHY / 0700.HK):** Near-term profit compression from tripling capex; the payoff depends on HY4 competitiveness. A longer-duration AI infrastructure bet within China.
- **Memory suppliers (Samsung, SK Hynix, Micron / MU):** The recurring “memory chip costs soaring” theme points to a broad-based memory up-cycle with pricing power — an indirect but consistent read-through.
- **HubSpot (HUBS):** Acquirer of Warmly; incremental exposure to the agentic-AI/data-value theme.
- **OpenAI & Anthropic (private):** The AlphaSense efficiency findings support their premium-pricing thesis; Anthropic’s reported ~\$6B Decart AI talks signal continued aggressive vertical integration into efficiency/world-model tech.
- **Private/thematic:** Kalshi (prediction markets, pre-IPO, regulatory risk), SpaceX (compute reselling upside), and model-routing infrastructure (OpenRouter) are names to watch as the efficiency narrative gains traction.

Prepared automatically for The Information Daily. Coverage window: ~24 hours to 14 Aug 2026, 07:00 GMT+8. Sources: The Information (The Briefing, The Information AM, Applied AI, AI Agenda, Big Read) and ESPN as cited. Feature-article links may require a subscriber login.