

The Information Daily Briefing - 2026-10-01

Coverage window (backfill): The Information newsletters received on 1 October 2026, GMT+8: The Information AM (19:06), the SpaceX AI exclusive (21:01) and AI Agenda (22:01). No emails from The Information arrived between the 30 September briefing and 19:06 GMT+8 on 1 October. This backfill was produced on 2 October, and its stories are also covered, with later updates, in the 2026-10-02 briefing. Not investment advice.

Top stories

1. Micron's revenue nearly quintuples as HBM stays sold out

- **What:** Micron's revenue for the quarter ended 3 September rose almost 5x to **\$54.2bn**. Gross margin nearly doubled to **86.8%** and free cash flow reached **~\$33bn**, up from **~\$17.5bn** the quarter before. That is more than Apple's \$32bn in the July quarter.
- **Why:** Demand for high-bandwidth memory (HBM) for AI chips is outrunning supply. Price increases and long-term agreements (LTAs) with cloud providers have made the memory makers dependable cash generators.
- **When:** Reported 30 September (US). HBM is effectively sold out through next year. The real test comes when the LTAs expire in a few years.
- **Which:** Micron (MU), SK Hynix and Samsung (all three invested in Anthropic's latest round). Nvidia, Meta and Apple bear the cost.
- **How:** Nvidia has tested reduced-memory versions of Rubin Ultra and its chip-system prices are rising ~17% partly because of memory. Meta and Apple have flagged price increases.
- **Link:** [Revenue Quintupled at AI Memory Maker Micron](#)

2. Meta's Muse passes 3m weekly users

- **What:** Muse has **>3m weekly users who prompt it, >4m total weekly users** and **>1m prompting DAU**. A week after launch it had ~500k weekly users and ~250k DAU.
- **Why:** This is the clearest data yet on Meta's agentic AI push. It is still a small fraction of Meta's billions of users.
- **When:** Internal data reported 1 October. Muse launched 8 September.
- **Which:** Meta (META). Rivals include SpaceXAI's Grok Bot, OpenAI's Dots (launched at DevDay) and the startup Instinct.
- **How:** Most prompts come through the app rather than the web. The gap between the 4m and 3m figures is people who review or approve actions without giving Muse a new task.
- **Link:** [Muse Tops 3 Million Weekly Users](#)

3. Google launches Gemini 4 Argon at a large discount to rivals

- **What:** Google launched **Gemini 4 Argon**, its first frontier flagship since Gemini 3, for cyber defence, coding, legal and finance work. Introductory pricing is **\$2 per million input tokens and \$10 per million output tokens**, doubling later. Even the doubled rate is well below GPT-6 Astra and Claude Fable 5.1.
- **Why:** Google needs to show it is still in the frontier race with Anthropic and OpenAI.
- **When:** Announced 30 September (US). It goes first to Fairwind Program cyber partners, and broader release follows once safeguards are ready (no timeline given).
- **Which:** Alphabet (GOOGL), OpenAI, Anthropic.
- **How:** The release is gated and the pricing aggressive, adding to the frontier price war.
- **Link:** [Gemini 4 Argon](#)

4. Banks pull back from data-centre loan deals

- **What:** **Société Générale, SMBC and MUFG** are becoming more selective in data-centre lending.
- **Why:** Market volatility and rising yields are pushing lenders toward safer deals.
- **When:** Reported 1 October.
- **Which:** These banks previously led Oracle/OpenAI financings: SocGen led \$7.1bn for the first Stargate site, MUFG co-led \$38bn for two Oracle projects, and SMBC co-led \$18bn for Oracle's New Mexico project. Oracle (ORCL) and neoclouds are affected.
- **How:** Fewer banks are available to syndicate multibillion-dollar loans. Private credit is filling the gap; for example, PaleBlueDot AI is seeking \$600m from lenders including Brookfield for a Korean site serving Xiaohongshu.
- **Link:** [Banks Pull Back From Data Center Loan Deals](#)

5. SpaceX's AI unit becomes a neocloud and courted Microsoft

- **What:** SpaceX's AI unit held talks over the summer about leasing computing capacity to **Microsoft**. The current status is unclear.
- **Why:** Selling compute has become one of SpaceX's most lucrative businesses. Musk has embraced it after initial reluctance.
- **When:** Summer talks, reported 1 October.
- **Which:** SpaceX/SpaceXAI. Anthropic and Google are existing customers. Microsoft (MSFT) is a prospect.
- **How:** Compute-constrained labs are committing **billions of dollars a month**, putting SpaceX in competition with CoreWeave, Nscale and the hyperscalers.
- **Link:** [How SpaceX's AI Unit Turned Itself Into an AI Cloud Firm \(email\)](#)

6. China: DeepSeek moves onto Huawei, OpenAI accuses Moonshot, Tencent rents Oracle GPUs

- **What:** DeepSeek is open-sourcing tooling for **Huawei** AI chips (a TileLang port plus compute and communication

libraries). OpenAI says it disrupted a **model-distillation campaign** linked to Moonshot AI, which peaked at 16,000 requests from 4,000+ users on 24–25 July. **Tencent** signed a ~\$7bn, five-year lease with **Oracle** for ~100k advanced chips in Southeast Asia, with ~30% paid upfront.

- **Why:** Export controls are pushing China toward domestic chips and offshore rental capacity.
- **When:** Huawei training chips could reach DeepSeek as early as Q4 2026.
- **Which:** Huawei, DeepSeek, Moonshot, Tencent (0700.HK), Oracle (ORCL), Nvidia (NVDA).
- **How:** Huawei “supernode” systems were co-developed with DeepSeek. The Oracle lease routes Chinese demand through non-China data centres.
- **Links:** [DeepSeek and Huawei](#) · [OpenAI Accuses Moonshot](#)

7. Microsoft Science President Peter Lee steps down

- **What:** Peter Lee, who led Microsoft Research and later became President of Microsoft Science, is leaving the company.
- **Why:** He is the latest senior departure this year, after Rajesh Jha, Phil Spencer, Frank Shaw and Charlie Bell.
- **When:** Announced 30 September (US).
- **Which:** Microsoft (MSFT).
- **How:** Igor Tsyganskiy has run Microsoft Research since February.
- **Link:** [Peter Lee to Step Down](#)

8. Coding agents are pushing Vercel past \$600m ARR, plus private-market deals

- **What:** Vercel’s annualised revenue is **\$600m**, up 148% year on year. Coding agents drive about half of new business, versus less than 3% in January. It has 485k paying customers and intends to IPO (no timeline). Other deals: an **ElevenLabs** tender at **\$22bn**, **Novig** at \$2bn, CScale \$145m (Series C), Metaview \$60m and Flow Engineering at \$750m.
- **Why:** Agents default to Next.js and Vercel, which favours usage-priced, high-margin developer infrastructure.
- **When:** Reported 1 October.
- **Which:** Vercel (private). Competitors include Cloudflare (NET), Netlify, Amazon (AMZN) and OpenRouter (being acquired by Stripe).
- **How:** Usage-based pricing is replacing seats. Vercel’s AI SDK is downloaded ~33m times a week.
- **Link:** [Why Coding Agents Are Choosing Vercel \(email\)](#)

Also noted

- **Apple** will launch a smart-home hub on 13 October, alongside a new HomePod mini and Apple TV. ([Bloomberg](#))

Overall market trend, development and sentiment

Trend: Suppliers of scarce inputs (HBM, contracted compute) are posting record economics, while bank appetite for speculative data-centre debt is cooling. Frontier model prices continue to fall.

Developments: Micron’s quarter reset the memory bar. Google re-entered the frontier race on price. Consumer agents (Muse) are gaining traction. China is deepening its Huawei stack and renting offshore GPUs.

Sentiment: Cautiously constructive. Demand signals are strong, but credit is getting more selective.

Key questions being asked and how they were addressed

Question	How it was addressed
Is AI memory demand peaking?	No sign yet: HBM is sold out through next year under LTAs.
Can Google keep pace at the frontier?	Argon is a gated, low-price launch. Capability versus rivals is still unproven.
Is AI debt financing tightening?	Yes, at the margin. Big banks are more selective and private credit is stepping in.
Is Muse more than a launch spike?	Weekly prompting users rose ~6x in about three weeks. Monetisation is still open.

Key risks

Risk	Detail	Most exposed
Tighter credit	Smaller loan syndicates, higher cost of capital	ORCL, neoclouds
Memory cost squeeze	Hardware margins and pricing	NVDA, AAPL, META
Memory cyclicalit	LTAs lapse eventually	MU, SK Hynix, Samsung
China policy and IP	Distillation and offshore leasing	ORCL, NVDA
Frontier price war	API pricing compression	OpenAI, Anthropic

Potential investment implications (companies mentioned)

Company	Lean	Rationale
---------	------	-----------

Micron (MU)	Constructive	Record margin and FCF with LTA visibility. Watch for peak-cycle risk.
SK Hynix / Samsung	Constructive	Same HBM tailwind.
Nvidia (NVDA)	Neutral-to-constructive	Demand intact, but memory costs press on pricing and margins.
Meta (META)	Constructive	Muse usage is scaling quickly. Monetisation is the key watch item.
Alphabet (GOOGL)	Constructive	Back at the frontier with aggressive pricing.
Oracle (ORCL)	Cautious	Tencent lease adds backlog, but lender pullback and China policy risk weigh.
Microsoft (MSFT)	Neutral	Executive departures, and a potential buyer of SpaceX compute.
Cloudflare (NET)	Neutral-to-constructive	Agent-driven hosting demand is growing.
Tencent (0700.HK)	Neutral	Secured offshore compute, with US policy risk.

For information only. Not investment advice.